

A Computational-Experimental Framework for Rational Polymer Selection and Formulation Development of Spray-Dried Amorphous Solid Dispersions

Demonstrated with Indomethacin as a BCS Class II Model Drug

Master Research Framework — Version 2.0 (Frozen)

Revised by the Scientific Revision Committee in response to the Critical Scientific Review (CSR-2026-IND-ASD-001). This document supersedes Version 1.0 as the official research framework.

Interdisciplinary Computational Pharmaceutics Team

2026

Scientific Revision Committee

Editor-in-Chief · Associate Editor · Pharmaceutical Scientist · Computational Pharmaceutics Scientist · Formulation Scientist · QbD Expert · Industrial R&D; Scientist · Regulatory Scientist · Senior Biostatistician

Date frozen: August 2026

Table of Contents

Executive Summary

1. Introduction and Rationale
 2. Background and Literature Foundation
 3. Research Question and Hypotheses (Revised)
 4. Computational Architecture (Eight Layers, with Layer 5 Revision)
 5. End-to-End Computational Workflow (Revised)
 6. Mathematical Framework (Retained with Eq. 10 Revision and New Eqs. 11-12)
 7. Composite Compatibility Index (Revised)
 8. MCDA Framework Selection (Retained with Multi-Expert Protocol)
 9. Sensitivity Analysis (Revised)
 10. Uncertainty Quantification (Revised)
 11. Validation Strategy (Substantially Revised)
 12. Decision Framework Output Specification (Revised)
 13. Scientific Figures Catalogue (Revised)
 14. Python Implementation Plan (Retained with Module Additions)
 15. Statistical Framework (Substantially Revised)
 16. Experimental Framework (New in V2.0)
 17. Worked Example: Indomethacin + 6 Polymers (Revised)
 18. Failure Boundary Mapping (New in V2.0)
 19. Reviewer Simulation (Revised)
 20. Industrial Perspective (Revised)
 21. Final Framework Synthesis (Revised)
 22. Risk Register (New in V2.0)
 23. Implementation Roadmap (Revised)
 24. References (Updated)
- Appendix A — Review Matrix
- Appendix B — Change Log
- Appendix C — Rescored Quality Audit
- Appendix D — Freeze Statement

Executive Summary

This document presents Version 2.0 of the Master Research Framework for the rational polymer selection and QbD-driven formulation development of spray-dried amorphous solid dispersions (SD-ASDs) of indomethacin. Version 2.0 supersedes Version 1.0 as the official research framework and incorporates the mandatory and recommended revisions issued by the Scientific Revision Committee following a comprehensive rejection-oriented audit (Critical Scientific Review, CSR-2026-IND-ASD-001). The committee, comprising an Editor-in-Chief, Associate Editor, Pharmaceutical Scientist, Computational Pharmaceutics Scientist, Formulation Scientist, QbD Expert, Industrial R&D; Scientist, Regulatory Scientist, and Senior Biostatistician, classified every reviewer criticism as Mandatory, Recommended, Optional, Future Work, or Unsupported Reviewer Opinion and implemented every scientifically justified revision without exceeding the practical constraints of an M.Pharm project.

Version 1.0 was a research plan, not a completed study. The Critical Scientific Review correctly identified that no empirical data, no computational outputs, and no statistical analyses existed at the time of audit, and that submitting the framework as-is to a Q1 research-article journal would risk desk rejection within seven days. The committee accepts this assessment. Version 2.0 therefore reframes the framework explicitly as a methodologically frozen protocol awaiting experimental execution, and revises all novelty, hypothesis, validation, statistical, experimental, industrial, and publication claims to reflect this honest positioning. The composite quality score moves from the Version 1.0 self-assigned 78/100 to a committee-validated 69/100 in Version 2.0; this apparent decrease reflects the removal of inflated claims rather than a degradation of scientific substance, and the resulting document is more defensible under peer review.

Eight mandatory revisions have been implemented in Version 2.0. First, the experimental programme is now specified as a prerequisite for Q1 submission rather than a parallel aspiration, with a 40-week implementation roadmap ending at manuscript submission. Second, pre-registered statistical decision rules (primary endpoint, test statistic, null and alternative hypotheses, significance level, minimum effect size of interest, stopping rule, and borderline interpretation) are specified for each of the four retained hypotheses. Third, the single-batch pseudoreplication mitigation is replaced by a requirement for three independent spray-drying batches for the rank-1 polymer, with a documented sensitivity analysis as a fallback. Fourth, PXRD characterisation is extended to all top-2 ASDs at the initial timepoint, with T=end added for any ASD selected for stability. Fifth, 12-month long-term stability is added as planned future work with explicit acknowledgment that six-month accelerated data alone are insufficient for Q1 publication. Sixth, the Failure Boundary Map is redefined via logistic regression with a pre-registered boundary location, replacing the operational 80% dissolution threshold. Seventh, multi-expert AHP elicitation (three to five pharmaceutical scientists, geometric-mean aggregation) is now the primary method, with single-expert as fallback only. Eighth, a held-out test set and negative controls (polymers the screen rejected) are added to break the circularity of in-sample validation.

Five optional improvements have been implemented where feasible within M.Pharm constraints: baseline comparison against HSP-only and equal-weight averaging to demonstrate the marginal value of the full architecture; Bayesian sensitivity analysis for small-sample stability data; multi-pH dissolution (1.2, 4.5, 6.8); residual solvent testing by GC-HS per ICH Q3C; and particle-size plus morphology characterisation by SEM and laser diffraction. Three items are explicitly deferred to future work beyond the M.Pharm scope: full 12-month long-term stability, scale-up bridge studies from the Buchi B-290 to pilot-scale equipment, and multi-drug generalisation across five to ten BCS Class II compounds.

The novelty stance has been substantially revised. Three of the six novelty claims in Version 1.0 have been removed: the integrated eight-layer architecture (reclassified as a workflow description, not a scientific

contribution), the approximately 85% material reduction claim (reclassified as an arithmetic projection and relegated to supplementary calculation), and the framework transferability claim (reclassified as a hypothesis for future work, not a demonstrated novelty). The three retained novelty claims are the Composite Compatibility Index as a transparent integration device (with the explicit caveat that the HSP-Flory-Huggins cross-correlation requires PCA pre-processing), the operationalisation of failure boundary mapping for indomethacin SD-ASDs (bounded wording, acknowledging prior design-space concepts by Yu and Scheirs), and the open-source Python and R release with Docker containerisation. No claim of a new physical model, equation of state, or thermodynamic theory is made.

The hypothesis architecture has been revised to remove the cascading-failure fragility identified by the reviewer. H1 (computational rank predicts experimental rank, Spearman rho greater than or equal to 0.70) and H2 (Gordon-Taylor Tg RMSE less than or equal to 10 K) are retained as primary hypotheses but H1 is explicitly marked as exploratory at n equals 6 to 8 polymers and confirmatory only at n greater than or equal to 20. H3 (rank-1 polymer shows higher dissolution than rank-2) is reclassified as secondary. H4 (formulation remains within the failure boundary through six-month stability) is reclassified as secondary with an empirical fallback boundary independent of computational predictions. H5 (transferability to other BCS Class II drugs) is dropped as a formal hypothesis and reframed as a Discussion topic on generalisability, because a literature-anchored validation cannot falsify the framework given uncontrolled published data. A contingency plan is specified in the event that H1 fails: the manuscript is restructured as an exploratory computational case study with a methodological protocol contribution.

The publication target has been revised to reflect scope rather than impact factor. Version 1.0 named Molecular Pharmaceutics (impact factor approximately 4.6) as the primary target. The committee endorses the reviewer's assessment that this scope is mismatched with a QbD-driven framework paper and recommends AAPS PharmSciTech as the primary submission target (scope fit for QbD and formulation work, acceptance probability 35 to 45%), with International Journal of Pharmaceutics as a strong secondary (formulation-focused). If the experimental programme cannot be completed within the M.Pharm timeline, the framework is reframed as a protocol paper for submission to MethodsX or BMJ Open (acceptance probability 60 to 75%). The graphical abstract specification has been updated to a minimum of 800 by 800 pixels, the figure count reduced from ten to eight for the main text with the remainder moved to supplementary material, and a data availability statement template has been added.

Technology Readiness Level remains at TRL 4 (laboratory validation), Industrial Readiness Level remains at 2 to 3 (validated in a laboratory setting but not in an industrial pipeline), and Scientific Readiness Level has been revised downward from 7 to 8 to 5 to 6 to honestly reflect the validation gaps. The framework is academically defensible as a methodological contribution but is not industrial-ready; bridging to TRL 6 would require pilot-scale spray drying, GMP batch documentation, computer system validation, a cost-of-goods model, freedom-to-operate analysis, and comparison to in-silico tools already used by industry (COSMO-RS, GastroPlus, Simcyp) — all out of M.Pharm scope and deferred to PhD or industrial collaboration.

Version 2.0 is hereby frozen. No further methodological changes shall occur before computational implementation begins. Only corrections arising from experimental evidence may modify the framework thereafter. The Change Log in Appendix B provides per-revision traceability from Version 1.0 to Version 2.0, the Review Matrix in Appendix A maps every reviewer criticism to its disposition, the Rescored Quality Audit in Appendix C documents the scoring changes, and the Freeze Statement in Appendix D formally concludes the revision.

1. Introduction and Rationale

1.1 The BCS Class II Problem

Approximately 70 to 80 percent of new chemical entities entering pharmaceutical development exhibit poor aqueous solubility, with the majority classified as Biopharmaceutics Classification System (BCS) Class II compounds characterised by low solubility and high permeability. For such compounds, oral bioavailability is dissolution-rate limited or solubility-limited, and conventional crystalline formulations frequently fail to deliver therapeutically effective plasma concentrations. The amorphous state, in which the molecule lacks long-range crystallographic order, offers a thermodynamic excess of free energy that translates into apparent solubility enhancements of one to several orders of magnitude relative to the stable crystalline form. However, the amorphous state is thermodynamically unstable, and recrystallisation during processing, storage, or dissolution can nullify the solubility advantage and introduce uncontrolled variability into the drug product.

Amorphous solid dispersions (ASDs) address this thermodynamic liability by molecularly dispersing the drug within a polymeric carrier that kinetically traps the amorphous state through a combination of anti-plasticisation (T_g elevation), specific intermolecular interactions (hydrogen bonding, ionic interactions), and reduced molecular mobility. Spray drying is among the most widely adopted manufacturing platforms for ASDs because it produces homogeneous, single-phase powders at scales compatible with early development and is readily transferable to commercial manufacturing. Despite these advantages, the rational selection of the polymer carrier remains an open problem and is the principal bottleneck addressed by this framework.

1.2 The Polymer Selection Bottleneck

Polymer selection in ASD development is currently dominated by empirical screening cascades that consume milligram-to-gram quantities of drug substance, weeks of analyst time, and substantial instrument access (DSC, XRPD, FTIR or Raman, dissolution apparatus, accelerated stability chambers). A typical screen evaluates five to fifteen polymers at two to four drug loadings, generating a matrix of 20 to 60 binary dispersions that must be prepared, characterised, and ranked. The rankings are then integrated informally by the formulation scientist using a mixture of quantitative data (T_g shift, single- T_g observation, absence of crystalline peaks) and qualitative judgement (handling, solvent compatibility, regulatory familiarity). This approach is expensive, poorly auditable, and difficult to reproduce across laboratories.

Computational pre-screening offers a complementary strategy in which drug-polymer compatibility is predicted from molecular and thermodynamic descriptors before any wet-lab work is initiated. When executed rigorously, a computational pre-screen can reduce the experimental matrix by an order of magnitude, focus the wet-lab effort on the most promising three to five candidates, and produce an explicit, reviewable record of why each polymer was retained or rejected. This record is directly aligned with the Quality by Design (QbD) expectations articulated in ICH Q8(R2), in particular the requirement to identify and control critical material attributes (CMAs) through mechanistic understanding rather than empirical testing alone.

It must be stated plainly, however, that a computational pre-screen is a planning tool and not a substitute for experimental validation. The framework presented in this document is a research plan that specifies how such validation will be conducted; it does not itself constitute validation. This distinction, drawn sharply by the Critical Scientific Review, is accepted by the committee and reflected throughout Version 2.0.

1.3 Scope and Objectives

This work designs a computational-experimental framework for rational polymer selection and prediction of formulation performance for spray-dried amorphous solid dispersions (SD-ASDs). The framework is demonstrated using indomethacin, a BCS Class II non-steroidal anti-inflammatory drug with well-characterised polymorphism, literature-reported miscibility with several pharmaceutical polymers, and established spray-drying processability. The framework itself is designed to be scientifically transferable to other BCS Class II drugs, but no claim of demonstrated generalisability is made; transferability is treated as a hypothesis for future work (see Section 3) rather than as a present-tense novelty claim.

Explicit limitations are stated where the framework is unlikely to transfer cleanly: ionic drugs where the neutral-state HSP approximation breaks down, polymorphs with dramatically different lattice energies, polymers with significant ionic character, and ternary systems involving surfactants or second polymers. Within these boundaries, the framework addresses the following eight objectives, revised from Version 1.0 to remove language implying that publication readiness is achievable without empirical validation.

The eight objectives are: (1) reduce experimental polymer screening by an order of magnitude relative to an unscreened 15 to 25 polymer library; (2) predict polymer suitability before wet-lab work begins, with explicit uncertainty quantification; (3) predict formulation risk (recrystallisation, phase separation, hygroscopicity) using a logistic-regression-based failure boundary map; (4) estimate formulation robustness via sensitivity and uncertainty analyses including Morris elementary effects and joint-distribution Monte Carlo; (5) support Quality by Design (ICH Q8, Q9, Q10) through transparent, auditable, multi-expert decision-making; (6) provide transparent, reproducible decisions with explicit assumptions, limitations, and pre-registered statistical decision rules; (7) be reproducible from a single open-source Python repository with Docker containerisation and Zenodo DOI archival; and (8) be of publication quality suitable for submission to a Q1 pharmaceuticals journal after empirical validation is completed.

1.4 Novelty Stance (Revised)

The novelty stance has been substantially revised in Version 2.0. The Critical Scientific Review identified that three of the six novelty claims in Version 1.0 were inflated, partially supported, or unfalsifiable as stated. The committee has implemented the reviewer's recommendation to remove these three claims from the novelty list and to reclassify them as workflow description, supplementary calculation, and future work respectively.

The retained novelty stance is as follows. Every equation, model, and decision-analysis method used in the framework is published in the peer-reviewed literature; no new physical model, equation of state, or thermodynamic theory is proposed. The scientific contributions of the framework are: (a) the Composite Compatibility Index (CCI) as a transparent integration device combining five normalised compatibility signals (HSP, Flory-Huggins χ , RDKit descriptors, Gordon-Taylor T_g , literature evidence) with AHP-derived weights, with the explicit caveat that the HSP and χ signals are cross-correlated (Spearman rho approximately 0.94 in the worked example) and therefore PCA pre-processing is mandatory before CCI computation; (b) the operationalisation of failure boundary mapping for indomethacin SD-ASDs via logistic regression, with bounded wording acknowledging that the design-space concept was defined by Yu (2008) and the failure-boundary concept was proposed by Scheirs; and (c) the open-source Python and R release with Docker containerisation and Zenodo DOI archival as the first reproducible computational-experimental framework of this scope for ASD polymer screening.

Removed from the novelty list are: the integrated eight-layer architecture (now treated as a workflow description in Section 4, not a scientific contribution, because prior work by Mistry and co-workers and by Nair and co-workers covers 10 to 12 of the 13 architectural components individually); the approximately 85%

material reduction claim (now treated as an arithmetic projection in supplementary material, because no prospective empirical 25-polymer screen was performed to validate the comparison); and the framework transferability claim (now treated as a hypothesis for future work in Section 3, because single-drug validation cannot demonstrate generalisability).

The committee endorses the reviewer's principle that novelty is necessary but not sufficient for Q1 acceptance, and that novelty is not the same as first application. The retained claims are defensible under peer review because they are bounded, honestly acknowledge prior art, and are accompanied by explicit validation pathways. The removed claims are not defensible because they either inflate a workflow description into a contribution, present an arithmetic projection as a finding, or treat a hypothesis as a demonstrated result.

2. Background and Literature Foundation

2.1 Hansen Solubility Parameters

Hansen Solubility Parameters (HSP) extend the Hildebrand solubility parameter by decomposing the cohesive energy density into three contributions: dispersion (δ_D), polar (δ_P), and hydrogen-bonding (δ_H). For two substances to be mutually soluble, their three-dimensional HSP coordinates must lie within a solubility sphere of radius R_o centred on the solute. The relative energy difference (RED) number, defined as R_a divided by R_o , provides a dimensionless compatibility metric: RED less than 1 indicates good affinity, RED equal to 1 indicates marginal affinity, and RED greater than 1 indicates poor affinity. HSP values may be obtained experimentally (Hansen, 2007) or estimated via group-contribution methods such as Hoftyzer-Van Krevelen or Stefanis-Panayiotou.

In ASD science, HSP has been used to predict drug-polymer miscibility with varying degrees of success. Greenhalgh and co-workers (1999) demonstrated a qualitative correlation between HSP distance and experimentally observed miscibility for several drug-polymer pairs. Thakral and co-workers (2012, 2020) extended this work and emphasised the importance of accurate HSP estimation. HSP-based screening is computationally cheap (seconds per pair on commodity hardware) but its predictive power is limited because it captures only cohesive energy considerations and does not account for specific intermolecular interactions such as ionic interactions or strong directional hydrogen bonds. This limitation is particularly relevant for indomethacin, which has two hydrogen-bond donors and four hydrogen-bond acceptors, and for which HSP-based predictions must be interpreted with explicit uncertainty.

2.2 Flory-Huggins Interaction Parameter

The Flory-Huggins (FH) theory provides a lattice-based thermodynamic description of polymer solutions and blends. The interaction parameter χ quantifies the enthalpic contribution to mixing and is the principal predictor of drug-polymer miscibility in the framework. The conventional screening threshold χ less than or equal to 0.5 is widely used as a necessary (though not sufficient) condition for miscibility, and the critical χ for phase separation is given by χ_c equals 0.5 times the quantity $(1 + 1/\sqrt{r_1} + 1/\sqrt{r_2})^2$.

The Lindvig conversion (Lindvig and co-workers, 2002) provides an empirical bridge from HSP distances to χ values, with an estimated relative error of approximately 15 percent on polymer-solvent systems. The Critical Scientific Review correctly notes that the Lindvig conversion was derived for polymer-solvent systems and may transfer imperfectly to drug-polymer systems, particularly for hydrogen-bonded drugs such as indomethacin. Version 2.0 acknowledges this limitation explicitly: the χ uncertainty for drug-polymer systems is increased to plus or minus 25 percent relative (revised from the Lindvig estimate of 15 percent), and

the framework recommends melting-point-depression (MPD) measurements for at least one drug-polymer pair as a calibration check on the Lindvig conversion. The MPD calibration is specified as a recommended addition in Section 16.

2.3 Gordon-Taylor Equation

The Gordon-Taylor (GT) equation predicts the glass transition temperature of a binary amorphous mixture from the glass transition temperatures and mass fractions of the pure components, weighted by a constant K . The Simha-Boyer rule (K equals the ratio of density-times-T_g products) is adopted in the framework because it requires only density and T_g, both of which are readily available. The conventional screening criterion requires that the T_g of the dispersion exceed the storage temperature by a sufficient margin, typically 50 K, to ensure kinetic stability against recrystallisation.

The Critical Scientific Review identified a systematic bias in the use of bulk crystalline density for indomethacin in the Simha-Boyer rule. The amorphous density of indomethacin is approximately 1.20 to 1.25 g/cm³, which is 5 to 10 percent lower than the crystalline density of 1.31 g/cm³ used in Version 1.0. This introduces a systematic bias in K of the same order, propagating to a T_g prediction error of approximately 3 to 5 K for indomethacin-PVP K30 at 30 percent drug loading. While this is within the RMSE acceptance threshold of 10 K, it is not negligible. Version 2.0 specifies that amorphous density should be measured where feasible and used in preference to crystalline density; where amorphous density is unavailable, the crystalline density is used with an explicit plus or minus 5 K systematic-bias flag on the resulting T_g prediction.

The optional Kwei correction (q times w_1 times w_2) partially addresses deviations from GT ideality for hydrogen-bonded systems, but the q -term must be fitted to experimental data. Version 2.0 specifies the Kwei correction as an optional refinement contingent on prospective T_g measurements for at least one drug-polymer pair, with the fitted q reported alongside the standard GT prediction in the Decision Report.

2.4 Molecular Descriptors via RDKit

RDKit-derived 2D molecular descriptors provide a complementary, structure-based signal that captures features not represented in the cohesive-energy-based HSP and chi scores. The descriptors used in the framework are molecular weight, topological polar surface area (TPSA), hydrogen-bond donor count (HBD), hydrogen-bond acceptor count (HBA), Crippen logP, rotatable bond count, aromatic ring count, and fractional polar surface area. For copolymers, a weighted average based on monomer mole fraction is used. These descriptors are computationally cheap, reproducible, and available from open-source software.

The descriptor-based compatibility score (s_{desc}) is a weighted combination of four sub-scores: HBD match (weight 0.3), HBA match (weight 0.3), TPSA proximity (weight 0.2), and aromatic ratio (weight 0.2). These sub-weights are specified explicitly in Version 2.0 in response to the reviewer's criticism that the descriptor score was opaque in Version 1.0. A one-at-a-time sensitivity analysis on these sub-weights is added to the standard sensitivity report.

2.5 Multi-Criteria Decision Analysis: AHP and TOPSIS

The Analytic Hierarchy Process (AHP; Saaty, 1980) elicits criterion weights from a pairwise comparison matrix using the principal eigenvector method, with a built-in consistency check (consistency ratio CR less than 0.10). The Technique for Order Preference by Similarity to Ideal Solution (TOPSIS; Hwang and Yoon, 1981) ranks alternatives by their Euclidean distance to the ideal and anti-ideal solutions, producing a closeness coefficient (CL) in the range 0 to 1. Both methods have mature open-source implementations in Python (pyAHP, scikit-criteria) and R (MCDM package).

Version 2.0 adopts multi-expert AHP elicitation as the primary method, in response to reviewer criticisms that single-expert elicitation is incompatible with ICH Q9 risk-management principles and with pharmaceutical decision governance. Three to five pharmaceutical scientists form the elicitation panel; geometric-mean aggregation produces the consolidated weight vector; inter-expert agreement is reported via Kendall's coefficient of concordance (W); and a documented revision protocol is invoked if CR exceeds 0.08 (tightened from the Saaty threshold of 0.10 in response to reviewer comment Q-MOD-001). Single-expert elicitation remains available as a fallback only when multi-expert elicitation is logistically infeasible, with an explicit flag in the Decision Report.

2.6 Quality by Design (ICH Q8, Q9, Q10)

ICH Q8(R2) defines Quality by Design (QbD) as a systematic approach to pharmaceutical development that begins with predefined objectives and emphasises product and process understanding and process control in an integrated manner. ICH Q9 addresses quality risk management and expects cross-functional input to risk decisions. ICH Q10 defines the Pharmaceutical Quality System (PQS) with four elements: change management, corrective and preventive action (CAPA), product quality review, and management responsibility.

The framework's alignment with ICH Q8 and Q9 is operational: the eight-layer architecture produces an auditable decision trail, the AHP-TOPSIS MCDA provides transparent weight derivation, and the multi-expert elicitation protocol satisfies the cross-functional input expectation of ICH Q9. Alignment with ICH Q10 is aspirational rather than operational: the PQS elements (change management, CAPA, PQR, management responsibility) are industrial-scale quality system components that are out of scope for an M.Pharm research framework but are acknowledged in Section 20 as future industrial-validation requirements. This honest distinction between operational and aspirational alignment is a revision from Version 1.0, which presented Q10 alignment without qualification.

2.7 Comparison to Existing In-Silico Tools (New in V2.0)

Pharmaceutical companies routinely use COSMO-RS, GastroPlus, Simcyp, and other in-silico tools for formulation and biopharmaceutical prediction. The framework presented here does not replace these tools and does not claim superiority over them. The framework's intended role is as an open-source, transparent pre-screen that narrows the polymer library before more expensive in-silico or wet-lab work is initiated. COSMO-RS provides quantum-chemistry-based predictions of drug-polymer miscibility with higher accuracy than the Lindvig conversion but requires commercial licensing and substantial computational resources (typically a Linux cluster with quantum-chemistry software). GastroPlus and Simcyp provide physiologically based pharmacokinetic (PBPK) modelling that is beyond the scope of this framework.

The committee endorses the reviewer's recommendation (criticism 9.7) that a comparative analysis with COSMO-RS would strengthen the framework's industrial positioning, but acknowledges that this comparison requires resources (COSMO-RS license, computational infrastructure) that are not available within the M.Pharm envelope. A COSMO-RS comparison is therefore specified as a future-work item in Section 21. The framework's value proposition in the present scope is that it offers a reproducible, open-source, zero-cost pre-screen that can be deployed in any academic or small-industry setting, complementing rather than competing with commercial in-silico tools.

3. Research Question and Hypotheses (Revised)

3.1 Primary Research Question

The primary research question addressed by the framework is restated from Version 1.0 with a bounded scope: can a computational polymer ranking based on HSP, Flory-Huggins χ , RDKit descriptors, Gordon-Taylor T_g , and AHP-TOPSIS multi-criteria decision analysis predict the experimentally observed miscibility and dissolution performance of spray-dried indomethacin ASDs, and can a logistic-regression-based failure boundary map identify the operating-condition region beyond which product quality fails? The question is bounded to indomethacin and to the six-polymer library specified in Section 17; generalisation to other BCS Class II drugs is treated as a hypothesis for future work (see Section 3.6) and not as a present-tense claim.

The Critical Scientific Review identified that Version 1.0 conflated ranking with dissolution performance, treating dissolution Q15 and Q30 as the sole basis for the experimental rank. Version 2.0 addresses this by specifying a composite experimental ranking criterion that combines T_g (single- T_g observation by DSC), miscibility (χ threshold and FTIR carbonyl shift), dissolution (multi-pH Q15 and Q30), and stability (six-month accelerated T_g retention and absence of crystalline peaks by PXRD). The composite criterion is computed as a weighted sum with weights pre-registered on Zenodo before wet-lab work begins.

3.2 Hypothesis Classification

The hypothesis architecture has been substantially revised to remove the cascading-failure fragility identified by the reviewer. Version 1.0 presented five hypotheses (H1 through H5) with implicit dependencies such that failure of H1 would propagate to H2, H3, and H4, collapsing four of five hypotheses simultaneously. Version 2.0 retains four hypotheses with explicit classification and contingency planning, and reclassifies the fifth (transferability) as a Discussion topic.

The four retained hypotheses are classified as primary (H1, H2), secondary (H3, H4), and exploratory (none retained as formal hypotheses; H5 is dropped and reframed as Discussion). Each hypothesis is paired with pre-registered statistical decision rules specifying the primary endpoint, test statistic, null and alternative hypotheses in formal notation, significance level (α equals 0.05), minimum effect size of interest, stopping rule, and borderline interpretation. The full pre-registration template is provided in Section 15.

3.3 H1 — Computational Rank Predicts Experimental Rank (Primary, Exploratory at $n=6-8$)

H1 (primary): the computational polymer ranking produced by the CCI-AHP-TOPSIS pipeline correlates with the experimental composite ranking (Section 3.1) with Spearman ρ greater than or equal to 0.70. Null hypothesis H0: ρ is less than 0.70. Alternative H1: ρ is greater than or equal to 0.70. Test statistic: Spearman rank correlation coefficient. Significance level: α equals 0.05. Minimum effect size of interest: ρ greater than or equal to 0.70. Stopping rule: data collection ceases after the pre-specified polymer set is characterised; no interim analysis. Borderline interpretation: if the 95 percent confidence interval for ρ includes 0.70, the result is recorded as 'exploratory, requiring confirmation at n greater than or equal to 20.'

The Critical Scientific Review correctly identified that at n equals 6 to 8 polymers, the 95 percent confidence interval for Spearman ρ at a point estimate of 0.83 spans approximately 0.20 to 0.98 (Fisher z -transformation), which is too wide to confirm or refute H1. Version 2.0 acknowledges this limitation explicitly: H1 is treated as exploratory at n equals 6 to 8 and confirmatory only at n greater than or equal to 20. A confirmatory computational expansion to n equals 20 polymers (literature-only, no wet-lab) is specified in the implementation roadmap (Section 23, Phase 2) as a mandatory step before any claim of H1 confirmation is made.

3.4 H2 — Gordon-Taylor Tg RMSE Within 10 K (Primary, Fully Answerable)

H2 (primary): the Gordon-Taylor predicted Tg of drug-polymer binaries agrees with the experimentally measured Tg (by DSC, single-Tg observation) with RMSE less than or equal to 10 K and MAE less than or equal to 8 K. Null hypothesis H0: RMSE is greater than or equal to 10 K. Alternative H1: RMSE is less than 10 K. Test statistic: RMSE and MAE with 95 percent bootstrap confidence intervals (n_boot equals 10,000). Significance level: alpha equals 0.05. Minimum effect size of interest: RMSE less than 10 K. Stopping rule: data collection ceases after the pre-specified polymer set is characterised. Borderline interpretation: if RMSE lies in the range 9 to 11 K, the result is recorded as 'marginally supported' and the Gordon-Taylor model is flagged for Kwei-correction fitting.

H2 is the only fully answerable hypothesis in the framework, as identified by the reviewer. The underlying Gordon-Taylor K via Simha-Boyer has a known systematic bias from the use of bulk crystalline density (Section 2.3); the framework specifies that amorphous density should be measured where feasible and used in preference, with the crystalline-density prediction reported alongside as a sensitivity check. The H2 acceptance criterion of RMSE less than or equal to 10 K is retained from Version 1.0.

3.5 H3 and H4 — Dissolution and Stability (Secondary)

H3 (secondary): the rank-1 polymer shows higher dissolution Q15 and Q30 than the rank-2 polymer at pH 6.8, with effect size Cohen's d greater than 0.5. Null hypothesis H0: Q15 and Q30 of rank-1 equals rank-2. Alternative H1: Q15 and Q30 of rank-1 greater than rank-2. Test statistic: two-sample t-test (or Mann-Whitney U if Shapiro-Wilk normality test rejects at p less than 0.05), with Cohen's d for effect size. Significance level: alpha equals 0.05. Minimum effect size of interest: Cohen's d greater than 0.5. The reviewer correctly identified that n equals 12 vessels per polymer is marginal for H3; Version 2.0 retains 12 vessels per polymer per pH medium (USP Apparatus II, 900 mL, 50 rpm, 37 plus or minus 0.5 degrees Celsius) and adds multi-pH dissolution (1.2, 4.5, 6.8) to strengthen the biopharmaceutical picture.

H4 (secondary): the formulation remains within the failure boundary (P(failure) less than 0.5 by logistic regression, Section 18) through six-month accelerated stability at 40 degrees Celsius and 75 percent relative humidity. Null hypothesis H0: P(failure) greater than or equal to 0.5 at any timepoint. Alternative H1: P(failure) less than 0.5 at all timepoints. Test statistic: logistic regression prediction at T equals 0, 1, 3, 6 months. Significance level: alpha equals 0.05. Borderline interpretation: if P(failure) lies in the range 0.30 to 0.70 at any timepoint, the formulation is classified as 'Warning' and the boundary is flagged for refinement. H4 depends on H1 through H3 in the sense that the failure boundary is defined using the computational ranking, but Version 2.0 specifies an empirical fallback boundary (based on observed CQA failure modes) that is independent of computational predictions, addressing the reviewer's criticism 3.1b.

3.6 H5 — Transferability (Dropped as Hypothesis, Reframed as Discussion)

H5 in Version 1.0 stated that the framework is transferable to other BCS Class II drugs, validated by literature-anchored computational screening of three to five additional drugs (felodipine, itraconazole, naproxen, carbamazepine, fenofibrate). The Critical Scientific Review correctly identified that this hypothesis is non-falsifiable within the framework's scope: published data for other drugs are not controlled (different polymers, different experimental conditions, different analytical methods), so a non-significant result would be attributable to data mismatch rather than framework failure. 'Transferable in principle' is a statement of belief, not a testable hypothesis.

Version 2.0 drops H5 as a formal hypothesis and reframes transferability as a Discussion topic on generalisability. The literature-anchored computational screening of three to five additional BCS Class II drugs

is retained as a supplementary analysis (not as hypothesis testing), with results presented descriptively and explicitly labelled as 'in-principle transferability, not demonstrated generalisation.' A full multi-drug wet-lab validation is specified as future work in Section 21, requiring a PhD-scale or industrial-collaboration programme.

3.7 Contingency Plan if H1 Fails

The Critical Scientific Review identified that Version 1.0 did not plan for the contingency that H1 fails. Version 2.0 specifies the following contingency: if H1 is not supported (95 percent confidence interval for ρ does not include 0.70, or point estimate is less than 0.50), the manuscript is restructured as an exploratory computational case study with a methodological protocol contribution, rather than as a confirmatory validation paper. The restructured manuscript would emphasise: (a) the framework's transparency and reproducibility as a methodological contribution; (b) the negative result itself as a scientifically informative finding (the CCI-AHP-TOPSIS pipeline does not predict experimental rank for indomethacin within the tested polymer library); (c) the diagnostic analysis of why H1 failed (PCA effective dimensionality, baseline comparison, Lindvig chi calibration); and (d) the protocol specification as a reusable tool for future researchers.

The contingency manuscript would target AAPS PharmSciTech or MethodsX rather than a Q1 research-article journal, with the framing 'this is a methodological protocol paper that reports a negative validation result and specifies the protocol for future confirmatory work.' This honest reframing avoids the temptation to present a failed validation as a partial success, which the reviewer correctly identified as a risk in Version 1.0's Success/Partial/Fail verdict structure.

4. Computational Architecture (Eight Layers, with Layer 5 Revision)

4.1 Architecture Overview

The framework is organised as eight functional layers, each with a well-defined input and output contract. The architecture enforces strict separation of concerns: each layer consumes the structured output of the layer below and produces a structured input for the layer above, with no implicit data flow across non-adjacent layers. This separation supports modularity (each layer can be replaced independently), testability (each function has unit tests), and reproducibility (the full pipeline runs from a single CLI command with a fixed random seed).

The eight layers are: Layer 1 Drug Knowledge, Layer 2 Polymer Knowledge, Layer 3 Descriptor Generation, Layer 4 Compatibility Prediction, Layer 5 Evidence Integration, Layer 6 Decision, Layer 7 Prediction, and Layer 8 Validation. The architecture itself is treated as a workflow description (not a novelty claim, per the revised novelty stance in Section 1.4). The scientific contributions reside in the specific methods used within Layers 5, 6, and 7, as detailed below.

4.2 Layers 1-4 (Retained from V1.0)

Layer 1 (Drug Knowledge) ingests canonical SMILES, pKa, Tm, Tg_drug (experimental preferred; Boyer-Beaman estimate as fallback), logP, density, BCS class, and known polymorphs. Plausibility validation: 300 K less than Tm less than 800 K; 0.8 g/cm³ less than density less than 2.0 g/cm³. For ionisable drugs such as indomethacin (pKa approximately 4.5), the layer tags the drug with its ionisation state and computes descriptors for both the neutral and ionised forms; HSP-based compatibility prediction is strictly valid only for the neutral-state interactions captured by the group-contribution method.

Layer 2 (Polymer Knowledge) stores for each polymer: monomer or repeat-unit SMILES, Mn, Mw, PDI, Tg, density, HSP (δ_D , δ_P , δ_H), functional-group inventory, regulatory status, and literature evidence

with the drug of interest. The default polymer library contains six polymers (PVP K30, PVP-VA 64, HPMCAS-L, Eudragit L100, Soluplus, HPMC E5) spanning neutral, anionic, and amphiphilic classes; the library is expandable to 30 to 50 polymers for industrial use (future work, Section 21).

Layer 3 (Descriptor Generation) produces 2D molecular descriptors via RDKit: MW, TPSA, HBD, HBA, logP (Crippen), rotatable bonds, aromatic ring count, fractional polar surface area, and HSP values via Hoftyzer-Van Krevelen group contribution. For copolymers, a weighted average based on monomer mole fraction is used.

Layer 4 (Compatibility Prediction) produces four normalised compatibility scores: s_{HSP} (from RED, with s_{HSP} equals $\max(0, 1 - \text{RED}/2)$), s_{chi} (from Lindvig chi, with s_{chi} equals $\max(0, 1 - \text{chi})$), s_{desc} (weighted combination of HBD match, HBA match, TPSA proximity, and aromatic ratio, each normalised to 0 to 1 within the candidate set), and s_{GT} (from Gordon-Taylor Tg, with s_{GT} equals $(Tg_{\text{mix}} - (Tg_{\text{drug}} + 30 \text{ K})) / (80 \text{ K} - 30 \text{ K})$, clipped to 0 to 1). A fifth score, s_{lit} , is the literature-evidence score (1.0 if miscibility evidence, 0.5 if no literature, 0.0 if immiscibility evidence). All five scores are normalised to the range 0 to 1 before integration.

4.3 Layer 5 — Evidence Integration (Revised: PCA Pre-Processing Mandatory)

Layer 5 combines the five normalised scores (s_{HSP} , s_{chi} , s_{desc} , s_{GT} , s_{lit}) into the Composite Compatibility Index (CCI). Version 1.0 used a direct weighted linear combination: CCI equals w_{HSP} times s_{HSP} plus w_{chi} times s_{chi} plus w_{desc} times s_{desc} plus w_{GT} times s_{GT} plus w_{lit} times s_{lit} . The Critical Scientific Review identified that three of the five inputs (HSP, chi, GT) are derived from cohesive energy density and are mutually correlated, with Spearman rho between HSP and chi equal to 0.94 in the worked example. The direct weighted combination therefore double-counts cohesive-energy compatibility, and the informational content of the CCI is substantially less than five independent signals.

Version 2.0 mandates Principal Component Analysis (PCA) pre-processing before CCI computation. The five scores are first assembled into an N-by-5 matrix (N equals number of candidate polymers) and PCA is applied after centering and scaling. The CCI is then computed as a weighted combination of the retained principal components (those explaining at least 95 percent of cumulative variance), with weights derived from AHP applied to the principal components rather than to the raw scores. The effective dimensionality of the CCI (number of retained PCs, percentage variance explained by PC1) is reported in the Decision Report.

If the first principal component explains more than 80 percent of the variance (likely given the HSP-chi correlation of 0.94), the CCI effectively reduces to a weighted first-PC score. This is reported transparently as 'the CCI in this candidate set is effectively one-dimensional; the five-signal integration provides redundancy rather than additional information.' The committee considers this honest reporting preferable to the Version 1.0 presentation, which implied five independent signals without verifying their independence.

Missing values are handled by imputation of a neutral score (0.5) and propagation of an explicit uncertainty flag, as in Version 1.0. The PCA is computed on the available data; if more than 20 percent of values are missing for a given polymer, that polymer is flagged and excluded from ranking.

4.4 Layer 6 — Decision (Revised: Multi-Expert AHP Primary)

Layer 6 applies AHP to derive criterion weights and TOPSIS to rank candidate polymers. Version 2.0 makes multi-expert AHP elicitation the primary method, in response to reviewer criticisms 4.2 and 9.4. The elicitation panel comprises three to five pharmaceutical scientists drawn from formulation science, analytical chemistry, regulatory affairs, and quality assurance, satisfying the ICH Q9 cross-functional input expectation. Each expert completes the pairwise comparison matrix independently; the consolidated weight vector is computed as the geometric mean of the individual weight vectors (geometric-mean aggregation, per Saaty's recommendation for group decision-making).

Inter-expert agreement is reported via Kendall's coefficient of concordance (W). If W is less than 0.50 (moderate agreement), the elicitation is repeated with a structured discussion session; if W remains below 0.50 after two rounds, the disagreement is reported as an explicit uncertainty in the Decision Report. The AHP consistency ratio threshold is tightened from 0.10 (Saaty convention) to 0.08 in response to reviewer comment Q-MOD-001. If CR exceeds 0.08 for any individual expert or for the aggregated matrix, the pairwise judgements are flagged and the expert is asked to revise; a documented revision protocol (re-elicitation of the most inconsistent pairwise comparisons, identified by the contribution to $\lambda_{\max}$) is invoked.

Single-expert elicitation remains available as a fallback only when multi-expert elicitation is logistically infeasible (for example, in a single-institution M.Pharm setting without access to a panel). In this case, the Decision Report carries an explicit flag 'single-expert elicitation; multi-expert confirmation recommended before wet-lab commitment.'

4.5 Layer 7 — Prediction (Revised: Logistic Regression FBM)

Layer 7 produces the formulation performance predictions: expected T_g with prediction interval, miscibility class (miscible, partially miscible, immiscible), stability tier (high, medium, low risk of recrystallisation over six months at 25 degrees Celsius and 60 percent RH, with humidity-conditional reporting at 40 degrees Celsius and 75 percent RH), risk profile (recrystallisation, phase separation, hygroscopicity), and the failure boundary map (FBM).

The FBM is substantially revised in Version 2.0 (see Section 18 for full specification). Version 1.0 defined the failure boundary operationally as dissolution Q30 less than 80 percent, which the reviewer correctly identified as a binary threshold rather than a statistical criterion. Version 2.0 replaces this with a logistic regression model: $\text{logit}(P(\text{failure}))$ equals β_0 plus β_1 times polymer_rank plus β_2 times inlet_T plus β_3 times drug_loading plus β_4 times $\text{feed_concentration}$, with the boundary defined as the $P(\text{failure})$ equals 0.5 contour. The Safe, Warning, and Failure regions are defined probabilistically: Safe is $P(\text{failure})$ less than 0.30, Warning is 0.30 less than or equal to $P(\text{failure})$ less than or equal to 0.70, and Failure is $P(\text{failure})$ greater than 0.70. The boundary location is pre-registered on Zenodo before wet-lab validation begins, addressing the reviewer's cherry-picking concern (criticism 8.6).

4.6 Layer 8 — Validation (Substantially Revised, see Section 11)

Layer 8 compares computational predictions against experimental data. The validation strategy is substantially revised in Section 11 to address the reviewer's criticisms of circular validation (criticism 7.1), in-sample validation (criticism 7.2), and the lack of baseline comparison (criticism 7.7). Key revisions include: (a) inclusion of one to two negative-control polymers (polymers the computational screen rejected, such as HPMC E5 or Eudragit L100) in the wet-lab validation to test the screen's specificity; (b) a held-out test set protocol in which AHP weights are derived from four polymers and tested on two unseen polymers; (c) leave-one-out cross-validation on the six-polymer worked example; (d) baseline comparison against HSP-only ranking and equal-weight averaging, with the full CCI required to outperform baselines by a pre-specified margin (delta Spearman rho greater than or equal to 0.10); and (e) explicit marking of validation results as exploratory at n equals 6 to 8 and confirmatory only at n greater than or equal to 20.

5. End-to-End Computational Workflow (Revised)

5.1 Step-by-Step Workflow Specification

The end-to-end computational workflow is deterministic: given the same inputs and the same random seed for the Monte Carlo analyses, the output is bit-for-bit identical across runs and machines. Reproducibility is enforced by version-pinning all dependencies, capturing the full input space (including AHP pairwise judgements from each expert) in version-controlled JSON configuration files, and seeding the Monte Carlo random number generator with a fixed seed (42). The workflow is invoked by a single CLI command (`python -m asd_mcda.cli --config config/workflow_config.yaml`) and completes in under five minutes on a standard laptop, including the 10,000-iteration Monte Carlo uncertainty quantification.

The eleven-step workflow is: (1) Input (drug JSON, polymer library CSV, AHP matrices from each expert); (2) Preprocessing (RDKit parsing, descriptor calculation); (3) Descriptor calculation (HSP via Hoftyzer-Van Krevelen, χ via Lindvig, T_g via Gordon-Taylor with Simha-Boyer K); (4) HSP scoring (s_{HSP} equals $\max(0, 1 - \text{RED}/2)$); (5) Flory-Huggins scoring (s_{χ} equals $\max(0, 1 - \chi)$); (6) Gordon-Taylor scoring (s_{GT} normalised between $T_{g_drug} + 30 \text{ K}$ and $T_{g_drug} + 80 \text{ K}$); (7) Evidence integration (PCA on five scores, then weighted combination into CCI using AHP weights applied to retained PCs); (8) Ranking (AHP weights from geometric-mean aggregation across experts, TOPSIS ranking with closeness coefficient CL); (9) Prediction (expected T_g with prediction interval, miscibility class, stability tier, FBM via logistic regression); (10) Validation (comparison against experimental data, baseline comparison, LOO cross-validation); (11) Output (Decision Report PDF and JSON, ranking CSV, sensitivity report CSV, uncertainty report CSV, validation report CSV).

5.2 Decision Gates (Revised)

Four decision gates are specified. Gate 1 (after HSP filter): if fewer than three polymers pass RED less than or equal to 1.0, the polymer library is flagged as too narrow and the framework halts with a request to expand the library. Gate 2 (after AHP consistency check): if CR exceeds 0.08 for any individual expert or for the aggregated matrix, the pairwise matrix is flagged and the most inconsistent judgements are revisited via the documented revision protocol; the framework does not proceed to ranking until CR is below 0.08 for all experts and the aggregated matrix.

Gate 3 (after validation): if RMSE exceeds 10 K, or Spearman rho is less than 0.70, or the full CCI does not outperform the HSP-only baseline by delta Spearman rho greater than or equal to 0.10, the framework enters a refinement loop. AHP weights are recalibrated based on the LOO cross-validation results, the polymer library is

reviewed for inclusion of additional candidates, and the Lindvig chi conversion is checked against MPD calibration data if available. Gate 3 failure triggers a documented refinement log entry and the framework is re-run with revised inputs.

Gate 4 (new in V2.0, after FBM estimation): the logistic regression FBM must satisfy two conditions before the boundary is reported as actionable. First, the boundary must be pre-registered on Zenodo before wet-lab validation begins (addressing the cherry-picking concern, criticism 8.6). Second, the boundary location must be reported with a bootstrap 95 percent confidence interval (n_{boot} equals 10,000) on the contour coordinates; if the CI width exceeds 20 percent of the design-space range in either dimension, the boundary is flagged as 'low confidence' and additional data points are requested before the boundary is used for go/no-go decisions.

6. Mathematical Framework (Retained with Eq. 10 Revision and New Eqs. 11-12)

6.1 Equations 1-9 (Retained from V1.0)

The nine core equations from Version 1.0 are retained without modification. Equation 1 (HSP distance, Hansen 2007): R_a equals the square root of (4 times ($\Delta_{D,A}$ minus $\Delta_{D,B}$) squared plus ($\Delta_{P,A}$ minus $\Delta_{P,B}$) squared plus ($\Delta_{H,A}$ minus $\Delta_{H,B}$) squared). Equation 2 (RED number): RED equals R_a divided by R_o . Equation 3 (Hofsteyer-Van Krevelen group contribution for HSP). Equation 4 (Flory-Huggins free energy of mixing per lattice site, Flory 1953 and Huggins 1942). Equation 5 (Lindvig conversion of HSP to chi, with empirical weights 0.6, 0.25, 0.25; Lindvig and co-workers 2002). Equation 6 (Gordon-Taylor Tg prediction, Gordon and Taylor 1952). Equation 7 (Simha-Boyer estimate of K, Simha and Boyer 1962). Equation 8 (AHP weight derivation from principal eigenvector, with consistency ratio CR equals CI divided by RI(n), where RI(5) equals 1.12; Saaty 1980). Equation 9 (TOPSIS closeness coefficient CL equals D_i minus divided by (D_i plus D_i minus)); Hwang and Yoon 1981). Full mathematical statements of these equations are provided in the notation appendix.

Supporting formulae retained from Version 1.0 include the critical chi for phase separation (χ_c equals 0.5 times (1 plus 1 over square root of r_1 plus 1 over square root of r_2) squared), the Boyer-Beaman Tg estimate (Tg_{drug} equals 0.7 times Tm_{drug}), the score normalisations, the justification trace decomposition (contribution_i equals w_i times s_i), and the decision confidence metric ($P(\text{top-1})$ equals the number of Monte Carlo runs in which polymer i is rank-1, divided by N).

6.2 Equation 10 — CCI (Revised: PCA Pre-Processing)

Equation 10 (revised): the Composite Compatibility Index is computed on the retained principal components of the five-score matrix rather than on the raw scores. Formally, let S be the N-by-5 matrix of normalised compatibility scores (s_{HSP} , s_{chi} , s_{desc} , s_{GT} , s_{lit}) for N candidate polymers. After centering and scaling, PCA produces the score matrix T (N-by-k) and loading matrix P (5-by-k), where k is the number of principal components retaining at least 95 percent of cumulative variance. The CCI for polymer i is then: CCI_i equals the sum over j from 1 to k of w_j times $T_{i,j}$, where w_j are the AHP-derived weights applied to the principal components (not to the raw scores).

The effective dimensionality k is reported alongside the CCI. If k equals 1 (PC1 explains at least 95 percent of variance), the CCI is effectively a weighted first-PC score and the five-signal integration provides redundancy rather than additional information; this is reported transparently. If k equals 2 or more, the integration provides non-redundant information from multiple components and the CCI is reported as a multi-dimensional integration.

The AHP weights w_j are derived from a separate pairwise comparison matrix in which the principal components (not the raw scores) are the criteria. Experts are asked to compare 'the principal component capturing cohesive-energy compatibility' (typically PC1) against 'the principal component capturing structure-based compatibility' (typically PC2, dominated by descriptors and literature) and so on. This re-framing of the AHP question is more meaningful than asking experts to compare the raw HSP and chi scores, which they cannot distinguish conceptually given their high correlation.

6.3 Equation 11 — Logistic Regression FBM (New in V2.0)

Equation 11 (new): the failure boundary is estimated by logistic regression. The binary outcome variable is 'operational failure' (defined as dissolution Q30 less than 80 percent at pH 6.8, OR Tg shift exceeding 5 K indicating phase separation, OR appearance of crystalline peaks by PXRD). The predictors are the four DoE factors: polymer rank (1, 2, 3, ...), inlet temperature (degrees Celsius), drug loading (percent w/w), and feed concentration (percent w/v). The model is: $\text{logit}(P(\text{failure})) = \beta_0 + \beta_1 \text{polymer_rank} + \beta_2 \text{inlet_T} + \beta_3 \text{drug_loading} + \beta_4 \text{feed_conc}$. The failure boundary is the contour where $P(\text{failure})$ equals 0.5.

The Safe, Warning, and Failure regions are defined probabilistically: Safe is $P(\text{failure})$ less than 0.30, Warning is $0.30 \leq P(\text{failure}) \leq 0.70$, and Failure is $P(\text{failure}) > 0.70$. This quantifies the Warning region (which was undefined in Version 1.0's 'approaches limit' language) and replaces the binary threshold with a continuous probability. The terminology uses 'operational failure' or 'CQA failure' rather than 'failure' alone, addressing the reviewer's criticism (8.5) that the latter overstates the consequences.

The logistic regression is fitted using maximum likelihood (scikit-learn LogisticRegression with L2 regularisation, C equals 1.0). Model fit is assessed by the area under the receiver operating characteristic curve (AUC-ROC); an AUC greater than or equal to 0.75 is required for the boundary to be reported as actionable. The minimum number of data points for boundary estimation is 20 to 30 (acknowledging the reviewer's criticism 8.3 that approximately 5 points is statistically indefensible); if fewer than 20 points are available, the boundary is reported as 'preliminary, requiring additional data.'

6.4 Equation 12 — Boundary Uncertainty (New in V2.0)

Equation 12 (new): the uncertainty in the failure boundary location is quantified via bootstrap. The logistic regression is refit on B equals 10,000 bootstrap resamples of the original data, producing a distribution of boundary locations. The 95 percent confidence interval on the boundary is the 2.5th and 97.5th percentiles of the bootstrap distribution. If the CI width exceeds 20 percent of the design-space range in either dimension, the boundary is flagged as 'low confidence' and additional data points are requested.

A Bayesian alternative (placing weakly-informative priors on the beta coefficients and computing the posterior on the boundary location via Markov Chain Monte Carlo) is specified as a recommended extension. The Bayesian approach is preferred for small samples because it provides better coverage and does not rely on asymptotic approximations. The committee endorses the reviewer's criticism (6.7) that the blanket exclusion of Bayesian methods in Version 1.0 was unjustified; Bayesian sensitivity analysis is therefore added to the statistical framework (Section 15) for stability data, and the Bayesian boundary estimation is specified as a recommended option for the FBM.

7. Composite Compatibility Index (Revised)

7.1 Scientific Justification

The Composite Compatibility Index (CCI) is a weighted linear combination of compatibility signals, designed as a transparent integration device that produces a single ranking score for each candidate polymer. The scientific justification for the CCI is that no single compatibility signal (HSP, χ , descriptors, Gordon-Taylor, literature) is sufficient to predict ASD performance: HSP captures only cohesive energy, χ captures only the enthalpic contribution to mixing, descriptors capture only structural features, Gordon-Taylor captures only the Tg-based kinetic stability, and literature captures only prior experimental evidence. A weighted combination is the simplest integration device that preserves additivity (and therefore auditability), is robust to missing values (via neutral-score imputation), and is mathematically consistent with the AHP-TOPSIS additive utility model.

The CCI as a transparency device is the retained novelty claim N2 of the framework (Section 1.4). The committee acknowledges the reviewer's caveat (criticism 2.2) that 'novel' and 'valid' are different concepts: the CCI may be novel and still be wrong if the HSP- χ cross-correlation (ρ equals 0.94) causes double-counting of cohesive-energy compatibility. Version 2.0 addresses this by mandating PCA pre-processing (Section 4.3, Equation 10), which extracts orthogonal information from the correlated inputs before the CCI is computed.

7.2 Mathematical Definition (Revised)

The mathematical definition of the CCI is revised from Version 1.0 to incorporate PCA pre-processing. The five normalised scores (s_{HSP} , s_{χ} , s_{desc} , s_{GT} , s_{lit}) are assembled into the N-by-5 matrix S . After centering and scaling, PCA produces the score matrix T and loading matrix P . The CCI for polymer i is CCI_i equals the sum over j from 1 to k of w_j times $T_{i,j}$, where k is the number of retained PCs (those explaining at least 95 percent of cumulative variance) and w_j are AHP-derived weights applied to the principal components.

The score normalisations are retained from Version 1.0: s_{HSP} equals $\max(0, 1 - \text{RED}/2)$; s_{χ} equals $\max(0, 1 - \chi)$; s_{desc} equals the weighted combination of HBD match (0.3), HBA match (0.3), TPSA proximity (0.2), and aromatic ratio (0.2); s_{GT} equals $(Tg_{\text{mix}} - (Tg_{\text{drug}} + 30 \text{ K})) / (80 \text{ K} - 30 \text{ K})$, clipped to 0 to 1; s_{lit} equals 1.0 if published miscibility evidence exists, 0.5 if no literature is available, 0.0 if published immiscibility evidence exists.

7.3 Weight Determination via Multi-Expert AHP

The AHP weights w_j are derived from multi-expert elicitation (Section 4.4). Three to five pharmaceutical scientists independently complete the pairwise comparison matrix on the principal components (not the raw scores, per Section 6.2). The consolidated weight vector is the geometric mean of the individual weight vectors. Inter-expert agreement is reported via Kendall's coefficient of concordance (W); the elicitation is repeated if W is less than 0.50. The AHP consistency ratio threshold is 0.08 (tightened from 0.10).

The default AHP matrix structure for the principal components is provided in Appendix D. In the worked example (Section 17), if PC1 captures cohesive-energy compatibility (HSP + χ + GT loadings) and PC2 captures structure-based compatibility (descriptors + literature loadings), the default matrix gives PC1 moderate importance over PC2 (pairwise value 2), reflecting the greater mechanistic weight of cohesive-energy compatibility in ASD science. Experts may revise this judgement based on their domain expertise.

7.4 Advantages (Retained)

The CCI retains the advantages identified in Version 1.0: transparency (every weight and score is explicit and auditable), robustness to missing values (neutral-score imputation with uncertainty flag), mathematical consistency with AHP-TOPSIS additive utility, reproducibility (single-command execution with pinned dependencies and fixed seed), and sensitivity-friendliness (the additive structure supports OAT and Morris sensitivity analyses). The PCA pre-processing adds a sixth advantage: explicit handling of input correlation, which Version 1.0 lacked.

7.5 Limitations (Revised)

The CCI retains the following limitations from Version 1.0, with revisions. First, the CCI assumes linear additivity of compatibility contributions; non-linear interactions (synergistic hydrogen bonding and pi-pi stacking) are not captured. The PCA pre-processing mitigates this for the HSP-chi-GT cluster but does not address non-linearity between the cohesive-energy cluster and the descriptor-literature cluster. Second, AHP weight derivation is subjective; the multi-expert elicitation protocol reduces but does not eliminate this subjectivity. Third, the *s_{lit}* score is binary (1.0, 0.5, 0.0) and does not capture the strength of literature evidence; a graded literature score (based on the number and quality of supporting studies) is specified as a future refinement. Fourth, the CCI does not explicitly account for drug loading effects; the loading-range analysis (10 to 50 percent w/w in 5 percent steps, introduced in Version 1.0's reviewer-response section) is retained.

Fifth (revised): the cross-correlation between HSP, chi, and GT (all derived from cohesive energy density) is now explicitly addressed by PCA pre-processing. The committee acknowledges the reviewer's criticism (4.5) that the effective dimensionality of the CCI was unknown in Version 1.0; the PCA-based CCI reports the effective dimensionality transparently, and if PC1 explains more than 80 percent of variance, the CCI is reported as 'effectively one-dimensional.' Sixth (new): the PCA pre-processing itself introduces a limitation: the principal components are linear combinations of the original scores and may not have direct physical interpretation. The committee considers this an acceptable trade-off for the gained orthogonality and explicit dimensionality reporting.

7.6 Validation of the CCI (Revised)

The CCI is validated against experimental miscibility data using three protocols. First, the CCI ranking is compared against the qualitative literature ranking using Spearman rho and Kendall tau; a Spearman rho greater than or equal to 0.70 is considered acceptable evidence that the CCI captures the relevant physics, but is explicitly marked as exploratory at *n* equals 6 to 8 and confirmatory only at *n* greater than or equal to 20.

Second, leave-one-out cross-validation (LOO CV) is performed on the six-polymer worked example. Each polymer is held out in turn; the AHP weights are derived from the remaining five polymers; the held-out polymer is ranked; the rank is compared against the rank obtained when all six polymers are used. LOO CV tests the sensitivity of the ranking to individual polymer inclusion and provides a more honest estimate of generalisation error than in-sample validation.

Third, a held-out test set protocol is specified: AHP weights are derived from four polymers (PVP K30, PVP-VA 64, HPMCAS-L, Soluplus), and the framework is tested on two unseen polymers (Eudragit L100, HPMC E5). This protocol addresses the reviewer's criticism (4.7, 7.2) that the six-polymer worked example is in-sample. The held-out test set is small (*n* equals 2) but provides a genuine external validation check.

7.7 Effective Dimensionality Reporting (New in V2.0)

The framework mandates reporting of the PCA effective dimensionality alongside the CCI. The report includes: (a) the eigenvalues and percentage variance explained by each principal component; (b) the number of retained PCs (k) and the cumulative percentage variance explained; (c) the loadings of each original score on each retained PC, with interpretation (e.g., 'PC1 is dominated by HSP, chi, and GT loadings, representing cohesive-energy compatibility; PC2 is dominated by descriptors and literature, representing structure-based compatibility'); (d) the AHP weights applied to the retained PCs; and (e) the effective dimensionality classification (1D if k equals 1, multi-dimensional if k is greater than or equal to 2).

This reporting addresses the reviewer's criticism (4.5) that the informational content of the CCI was unknown in Version 1.0. If the effective dimensionality is 1, the committee recommends that the manuscript explicitly state 'the CCI in this candidate set is effectively one-dimensional, dominated by cohesive-energy compatibility; the descriptor and literature signals provide redundancy rather than additional information.' This honest reporting is more defensible than implying five independent signals were integrated.

8. MCDA Framework Selection (Retained with Multi-Expert Protocol)

8.1 Comparison of MCDA Methods (Retained)

Six MCDA methods were compared in Version 1.0: Weighted Sum, Weighted Product, AHP, TOPSIS, PROMETHEE, VIKOR, and ELECTRE. The comparison criteria were mathematical transparency, implementation ease, software availability, weight sensitivity, conflict handling, and M.Pharm feasibility. AHP was selected for weight derivation because of its transparent, auditable weight vector, built-in consistency check (CR less than 0.08 in V2.0), wide use in pharmaceutical formulation literature (Arslan and co-workers 2018; Tomer and co-workers 2020), and mature open-source implementations (pyAHP, scikit-criteria, MCDM package). TOPSIS was selected for ranking because of its mathematical simplicity (Euclidean distances), intuitive closeness coefficient (CL in 0 to 1), good handling of conflicting criteria, open-source availability, and pharmaceutical precedent (Behzadian and co-workers 2012 review of 265 TOPSIS applications).

The rejected methods remain rejected: Weighted Sum does not handle conflicting criteria well; Weighted Product is brittle to zero values; PROMETHEE, VIKOR, and ELECTRE have non-linear responses that complicate sensitivity analysis. The committee endorses this selection and makes no change in V2.0.

8.2 Justification of AHP-TOPSIS (Retained)

The justification for AHP-TOPSIS is retained from Version 1.0: AHP provides the weight vector with a consistency check, TOPSIS provides the ranking with an intuitive closeness coefficient, and both have mature open-source implementations. The combination is widely used in pharmaceutical formulation literature and is implementable within the M.Pharm envelope.

8.3 Recommended AHP Matrix Structure (Retained)

The default AHP matrix structure is provided in Appendix D. In Version 2.0, the matrix is applied to the retained principal components rather than to the raw scores (per Section 6.2). The default Saaty 1-9 scale is used, with reciprocals when criterion j is more important than criterion i . If CR exceeds 0.08, the most inconsistent pairwise comparisons are identified (by their contribution to $\lambda_{\max}$) and revisited.

8.4 Multi-Expert Elicitation Protocol (New in V2.0)

The multi-expert elicitation protocol is specified as follows. Step 1: assemble a panel of three to five pharmaceutical scientists representing formulation science, analytical chemistry, regulatory affairs, and quality assurance. Step 2: each expert independently completes the pairwise comparison matrix on the retained principal components, using the Saaty 1-9 scale. Step 3: compute the individual weight vector for each expert using the principal eigenvector method. Step 4: compute the consolidated weight vector as the geometric mean of the individual weight vectors (geometric-mean aggregation, per Saaty's recommendation for group decision-making). Step 5: compute the AHP consistency ratio (CR) for each individual expert and for the aggregated matrix; if any CR exceeds 0.08, invoke the documented revision protocol. Step 6: compute Kendall's coefficient of concordance (W) across the individual weight vectors; if W is less than 0.50, repeat the elicitation with a structured discussion session. Step 7: report the consolidated weight vector, individual weight vectors, CR values, and W in the Decision Report.

The documented revision protocol (invoked when CR exceeds 0.08) is: (a) identify the pairwise comparison with the largest contribution to $\lambda_{\max}$ (the most inconsistent judgement); (b) ask the expert to revisit that judgement, with a structured prompt explaining the inconsistency; (c) re-compute CR; (d) repeat until CR is below 0.08 or until three revision rounds have been completed, at which point the expert's matrix is flagged as 'requires supervisor review.'

If multi-expert elicitation is logistically infeasible (single-institution M.Pharm setting without access to a panel), single-expert elicitation is used as a fallback. In this case, the Decision Report carries an explicit flag 'single-expert elicitation; multi-expert confirmation recommended before wet-lab commitment' and the framework's confidence tier is downgraded by one level (e.g., from High to Moderate).

9. Sensitivity Analysis (Revised)

9.1 Designed Sensitivity Analyses

Three sensitivity analyses are specified. First, one-at-a-time (OAT) weight perturbation: each AHP weight is multiplied by 1.5 and 0.5 in turn (with renormalisation), and the ranking is recomputed. The OAT analysis identifies which individual weights have the largest effect on the ranking. Second, Monte Carlo weight variation: all weights are drawn from a uniform distribution centred on the AHP-derived values with plus or minus 20 percent bounds, normalised; N equals 10,000 iterations. The median Spearman rho across the 10,000 runs is reported, with a target of rho greater than or equal to 0.90 for stable ranking.

Third (new in V2.0): Morris elementary effects screening. The Morris method samples the weight space along r equals 10 trajectories, each of length k plus 1 (where k is the number of weights), producing r times (k plus 1) model evaluations. For each weight, the Morris method computes the mean (μ) and standard deviation (σ) of the elementary effects across trajectories. High μ indicates a weight with large main effect; high σ indicates a weight involved in interactions or non-linear effects. The Morris method detects interaction effects that the OAT analysis misses, addressing the reviewer's criticism (Q-MOD-002) that OAT cannot detect interactions.

The Morris screening is implemented using the SALib Python library (SALib.sample.morris and SALib.analyze.morris). The results are reported as a Morris scatter plot (μ on x-axis, σ on y-axis, one point per weight) in the sensitivity report. Weights with μ greater than 0.10 and σ greater than 0.05 are flagged as 'dominant and interactive' and are highlighted for expert review.

9.2 Essential vs Optional Analyses

The three sensitivity analyses (OAT, Monte Carlo, Morris) are all classified as essential. The threshold analysis (minimum weight perturbation required to change top-1 polymer) is retained from Version 1.0 as essential. The pairwise correlation analysis between the five compatibility scores is retained as informative but optional; in Version 2.0, this analysis is subsumed by the PCA pre-processing (Section 4.3), which provides a more rigorous treatment of input correlation.

9.3 Acceptance Criteria (Revised)

The acceptance criteria are revised from Version 1.0. First, the top-1 polymer must remain unchanged in at least 80 percent of Monte Carlo runs for the decision to be classified as 'robust'; this threshold is retained but the committee acknowledges the reviewer's question (4.2) of why 80 percent and not 95 percent. The 80 percent threshold is justified operationally: it corresponds to a less than 1-in-5 chance of a different top-1, which the committee considers an acceptable risk for a pre-screening tool whose output will be confirmed by wet-lab validation. A 95 percent threshold would be appropriate for a final go/no-go decision tool but is overly stringent for a pre-screen.

Second, the median Spearman rho across the 10,000 Monte Carlo runs must be at least 0.90 for the ranking to be classified as stable. Third (new), the Morris screening must not flag any weight as both dominant (μ greater than 0.10) and interactive (σ greater than 0.05) without explicit expert review; if a weight is so flagged, the expert panel revisits the corresponding pairwise comparison. Fourth, the threshold analysis: if the minimum weight perturbation required to change top-1 exceeds plus or minus 30 percent of the AHP-derived value, the top-1 is classified as robust.

If the top-1 polymer is unchanged in 50 to 80 percent of Monte Carlo runs, the decision is classified as 'moderately robust' and the framework proceeds to wet-lab with top-2 and top-3. If the top-1 is unchanged in less than 50 percent of runs, the decision is classified as 'fragile' and the framework revisits the AHP matrix, expands the polymer library, or both.

9.4 Sobol Indices (Removed from Primary Analysis)

The Critical Scientific Review correctly identified (criticism 6.8) that Sobol first-order sensitivity indices at N equals 8 polymers (88 model evaluations) cannot be reliably estimated. The minimum recommended N for Sobol is 100 to 1000; below this, estimates are noisy and confidence intervals are wide. The S_i greater than 0.20 threshold for 'dominant' inputs was arbitrary and the estimates at N equals 8 cannot support this classification.

Version 2.0 removes Sobol indices from the primary sensitivity analysis. The Morris elementary effects method (Section 9.1) provides a more appropriate screening for the small- N regime. If a future expansion to N greater than or equal to 100 polymers is undertaken (future work, Section 21), Sobol indices may be re-introduced with the SALib implementation and confidence intervals reported.

10. Uncertainty Quantification (Revised)

10.1 Sources of Uncertainty and Their Magnitudes

Every computational input to the framework carries uncertainty. The seven principal sources are: HSP group-contribution error (plus or minus 1.5 MPa^{0.5} per component, normal distribution, Stefanis and Panayiotou 2008); Lindvig chi conversion error (plus or minus 25 percent relative for drug-polymer systems, normal distribution, revised from plus or minus 15 percent in V1.0); logP estimation error (plus or minus 0.7 log units, normal distribution, RDKit Crippen); Tg_{drug} Boyer-Beaman estimate (plus or minus 10 K, normal distribution, Baird and co-workers 2010); Tg_{polymer} literature (plus or minus 3 K, normal distribution, Brandrup and co-workers 1999); density estimation error (plus or minus 0.05 g/cm³, normal distribution, Girolami 1994); and AHP weight subjectivity (plus or minus 20 percent relative, uniform distribution, Saaty 1980).

10.2 Joint Distribution Propagation (Revised)

Version 1.0 specified Monte Carlo with N equals 10,000 iterations and plus or minus 20 percent uniform perturbation on AHP weights, but did not specify the joint distribution of the seven uncertainty sources or the propagation method. The Critical Scientific Review correctly identified (criticism 4.6) that this is sampling, not uncertainty quantification. True UQ requires joint propagation of all seven sources.

Version 2.0 specifies the joint distribution as follows. The seven sources are sampled jointly from a multivariate normal distribution (for the six normally-distributed sources) combined with a uniform distribution (for the AHP weights). The default covariance structure is diagonal (independent sources), which is a simplification acknowledged explicitly: HSP and chi are correlated (both derive from cohesive energy density) and the diagonal covariance understates the joint uncertainty. A future refinement (specified as recommended) is to estimate the covariance from the HSP-chi correlation observed in the candidate set and propagate it through the multivariate normal.

The Monte Carlo procedure is: for each of N equals 10,000 iterations, sample the seven uncertainty sources from the joint distribution, recompute the five compatibility scores, recompute the CCI (including PCA re-fit), recompute the AHP-TOPSIS ranking, and record the top-1 polymer. The decision confidence metric P(top-1) is the fraction of iterations in which each polymer is rank-1. The 95 percent prediction interval for the predicted Tg is the 2.5th and 97.5th percentiles of the iteration-specific Tg predictions.

10.3 Decision Confidence Metric

The decision confidence metric P(top-1) is retained from Version 1.0: P(top-1) equals the number of Monte Carlo runs in which polymer i is rank-1, divided by N. The confidence tiers are retained: P(top-1) greater than or equal to 0.70 is high-confidence (proceed to wet-lab with top-1 only); 0.40 less than or equal to P(top-1) less than 0.70 is moderate-confidence (proceed with top-2 and top-3); P(top-1) less than 0.40 is low-confidence (revisit AHP matrix, expand library, or both).

10.4 Reporting Requirements

The uncertainty quantification report includes: (a) the input uncertainty magnitudes with citations; (b) the joint distribution specification (default diagonal covariance with explicit acknowledgment of the HSP-chi correlation simplification); (c) the Monte Carlo sample size (N equals 10,000); (d) the convergence check results (Section 10.5); (e) the CCI distributions per polymer (violin plot); (f) the P(top-1) bar chart with the 0.70 confidence threshold; (g) the chosen confidence threshold with justification.

10.5 Convergence Criterion (New in V2.0)

The Critical Scientific Review (Q-MOD-021) asked what convergence criterion is used for the Monte Carlo. Version 2.0 specifies two convergence criteria, either of which must be satisfied. First, the Gelman-Rubin R-hat statistic is computed across 5 independent Monte Carlo chains (each with N equals 10,000 iterations and different random seeds); R-hat less than 1.01 for the P(top-1) values indicates convergence. Second, the coefficient of variation of P(top-1) between the N equals 10,000 run and the N equals 50,000 run must be less than 5 percent for each polymer. If neither criterion is satisfied, the Monte Carlo is re-run with N equals 50,000 and the convergence check is repeated.

11. Validation Strategy (Substantially Revised)

11.0 Framework Development vs Validation vs Generalization

The Critical Scientific Review identified that Version 1.0 conflated three distinct stages: framework development (specifying the architecture, equations, and parameters), framework validation (testing the framework's predictions against experimental data), and framework generalization (extending the framework to other drugs, polymers, or processing conditions). Version 2.0 distinguishes these three stages explicitly.

Framework development is the activity of specifying the computational architecture (eight layers), the mathematical framework (Equations 1-12), the AHP matrix structure, the polymer library, and the implementation plan. This stage is complete when the framework can be executed end-to-end on a test case (the indomethacin plus six polymers worked example in Section 17). Framework validation is the activity of testing the framework's predictions against experimental data. This stage is complete when the validation metrics (RMSE, Spearman rho, top-1 agreement, baseline comparison) have been computed and the acceptance criteria have been met (or, if not met, the framework has been refined and re-validated). Framework generalization is the activity of extending the framework to other drugs, polymers, or processing conditions. This stage is complete when the framework has been validated on at least three BCS Class II drugs (multi-drug validation) and the validation metrics have been confirmed across drugs.

Version 2.0 is at the end of framework development and the beginning of framework validation. Framework generalization is specified as future work (Section 21). The committee emphasises that any claim of generalization before multi-drug validation is unsupported and must not be made in the manuscript.

11.1 Negative Controls (New in V2.0)

The Critical Scientific Review (criticism 7.1) identified that Version 1.0's validation is circular: the computational predictions determine which polymers proceed to wet-lab, and the wet-lab data are then used to validate the predictions. This circularity means the framework is never tested on polymers it rejected, so its specificity (ability to correctly reject unsuitable polymers) is never assessed.

Version 2.0 adds one to two negative-control polymers to the wet-lab validation. The negative controls are polymers that the computational screen rejected (e.g., HPMC E5 rank-5, Eudragit L100 rank-6 in the worked example). These polymers are taken forward to wet-lab characterisation (DSC, FTIR, PXRD, dissolution) alongside the top-2 computational selections. If the negative controls perform poorly in wet-lab (e.g., show crystalline peaks by PXRD, low Tg, or poor dissolution), this confirms the screen's specificity. If the negative controls perform well, this challenges the screen's specificity and triggers a refinement of the AHP weights or the polymer library.

The negative-control protocol breaks the circularity of in-sample validation by testing the framework's predictions on polymers it would have rejected. The cost is modest (one to two additional polymers through the wet-lab pipeline) and the scientific value is substantial (genuine assessment of specificity).

11.2 Held-Out Test Set (New in V2.0)

The Critical Scientific Review (criticism 4.7, 7.2) identified that the six-polymer worked example is in-sample: the AHP weights, polymer library, and computational pipeline were all derived from literature that includes these six polymers. No held-out test set was specified. Version 2.0 adds two protocols.

First, leave-one-out cross-validation (LOO CV) on the six-polymer set: each polymer is held out in turn; the AHP weights are derived from the remaining five polymers; the held-out polymer is ranked; the rank is compared against the rank obtained when all six polymers are used. LOO CV provides a more honest estimate of generalisation error than in-sample validation.

Second, a held-out test set protocol: AHP weights are derived from four polymers (PVP K30, PVP-VA 64, HPMCAS-L, Soluplus), and the framework is tested on two unseen polymers (Eudragit L100, HPMC E5). The held-out test set is small (n equals 2) but provides a genuine external validation check. The committee acknowledges that n equals 2 is too small for confirmatory validation; the purpose is to detect gross failures of the screen's specificity, not to confirm the framework.

11.3 Validation Metrics and Acceptance Criteria (Revised)

The validation metrics are retained from Version 1.0 with explicit exploratory/confirmatory classification. Continuous metrics (for Tg prediction): RMSE less than or equal to 10 K (acceptance criterion, but marked exploratory at n equals 6 to 8 and confirmatory at n greater than or equal to 20); MAE less than or equal to 8 K (same classification). Ordinal metrics (for ranking prediction): Spearman rho greater than or equal to 0.70 (exploratory at n equals 6 to 8; confirmatory at n greater than or equal to 20); Kendall tau greater than or equal to 0.60 (same). Top-k agreement: top-1 agreement required (primary criterion); top-3 agreement greater than or equal to 1/3 (operational).

The 95 percent confidence interval for Spearman rho at n equals 6 and rho equals 0.83 spans approximately 0.20 to 0.98 (Fisher z-transformation); this interval is too wide to confirm or refute H1 and is reported explicitly. The confirmatory threshold is n greater than or equal to 20 polymers, at which the 95 percent confidence interval for rho at rho equals 0.83 narrows to approximately 0.60 to 0.93.

11.4 Baseline Comparison (New in V2.0)

The Critical Scientific Review (criticism 7.7) identified that Version 1.0 does not compare the full eight-layer architecture against simpler baselines. Without baseline comparison, the framework cannot demonstrate that its complexity is justified. Version 2.0 adds two baseline comparisons.

First, HSP-only ranking: the polymers are ranked using s_{HSP} alone (the simplest compatibility signal), and the Spearman rho between the HSP-only ranking and the experimental ranking is computed. Second, equal-weight averaging: the five scores are combined with equal weights (0.20 each) without AHP or PCA, and the Spearman rho between this ranking and the experimental ranking is computed. The full CCI-AHP-TOPSIS pipeline must outperform both baselines by delta Spearman rho greater than or equal to 0.10 for the framework's complexity to be justified. If the full pipeline does not outperform the baselines, the committee recommends simplifying the framework (e.g., dropping the descriptor and literature scores if they do not contribute).

11.5 Equivalence Testing for FBM (New in V2.0)

The Critical Scientific Review (criticism 6.6) identified that Version 1.0's failure boundary definition (Q30 less than 80 percent) is a binary classification, not an equivalence test. The correct statistical framework is equivalence testing with a pre-specified margin. Version 2.0 adds the Two One-Sided Tests (TOST) procedure for the FBM boundary.

The TOST procedure is: at the boundary edge (where $P(\text{failure})$ equals 0.5 by logistic regression), the dissolution Q30 is measured experimentally and tested for equivalence to 80 percent using TOST with a pre-specified margin of plus or minus 5 percent. The null hypothesis is that Q30 is outside the equivalence interval (less than 75 percent or greater than 85 percent); the alternative is that Q30 is within the interval (75 to 85 percent). If the null is rejected (both one-sided p-values less than 0.05), the boundary is confirmed at that operating point. The TOST procedure is applied at each boundary point used for validation.

12. Decision Framework Output Specification (Revised)

12.1 Decision Report Template (Revised)

The Decision Report is generated as machine-readable JSON and human-readable PDF. The report fields are revised from Version 1.0 to include the new V2.0 outputs. The retained fields are: `selected_polymer`, `topsis_CL`, `confidence_P_top1`, `cci_value`, `predicted_Tg_K`, `Tg_prediction_interval`, `predicted_chi`, `miscibility_class`, `stability_tier`, `risk_recrystallization`, `risk_phase_separation`, `risk_hygroscopicity`, `justification_trace`, `sensitivity_summary`. The new fields are: `pca_effective_dimensionality` (number of retained PCs and percentage variance explained by PC1), `multi_expert_kendall_w` (inter-expert agreement), `baseline_comparison_delta` (delta Spearman rho vs HSP-only and equal-weight baselines), `fbm_boundary_logistic` (logistic regression beta coefficients and $P(\text{failure})$ equals 0.5 contour coordinates), `fbm_boundary_ci` (bootstrap 95 percent CI on contour coordinates), `held_out_test_results` (ranking and validation metrics on held-out polymers), `negative_control_results` (wet-lab performance of rejected polymers).

Table 12.1 — Decision Report Fields (V2.0)

Field	Type	Example	Source
selected_polymer	string	Soluplus	TOPSIS ranking
topsis_CL	float [0,1]	0.81	TOPSIS
confidence_P_top1	float [0,1]	0.78	Monte Carlo UQ
cci_value	float [0,1]	0.85	PCA + AHP
predicted_Tg_K	float	338	Gordon-Taylor
Tg_prediction_interval	[lo, hi]	[333, 343]	Monte Carlo
predicted_chi	float	0.28	Lindvig
miscibility_class	string	Miscible	chi vs chi_c
stability_tier	string	Medium-High (25C/60%RH); Medium-Low (40C/75%RH)	Humidity-conditional
pca_effective_dimensionality	object	k=2, PC1=72%, PC2=18%	PCA (NEW)
multi_expert_kendall_w	float [0,1]	0.62	Multi-expert AHP (NEW)
baseline_comparison_delta	object	vs HSP-only: +0.14; vs equal-weight: +0.09	Baseline comparison (NEW)
fbm_boundary_logistic	object	beta=[-2.1, 0.4, -0.02, ...]	Logistic regression (NEW)
fbm_boundary_ci	object	[[x_lo, x_hi, y_lo, y_hi]]	Bootstrap (NEW)
held_out_test_results	object	{polymer: rank_correct}	Held-out CV (NEW)
negative_control_results	object	{polymer: PXRD_crystalline=Yes}	Negative controls (NEW)

12.2 Justification Trace (Retained)

The justification trace decomposes the CCI into contributions from each retained principal component: contribution_j equals w_j times $T_{i,j}$. The trace is reported alongside the CCI value. For the worked example (Section 17), if PC1 (cohesive-energy compatibility) contributes 0.22 and PC2 (structure-based compatibility) contributes 0.08 to Soluplus's CCI of 0.30 (note: the V2.0 CCI is computed on PCs, not raw scores, so the numeric values differ from V1.0), the trace makes the relative contributions explicit.

12.3 Interpretation Guidance (Revised)

The interpretation guidance is revised to reflect the V2.0 confidence tiers and the exploratory status of n equals 6 to 8 validation. Example interpretation: 'Proceed to wet-lab screen of Soluplus at 30 percent w/w drug loading; predicted Tg equals 65 degrees Celsius (95 percent CI 60 to 70 degrees Celsius), with a 50 K margin above 25 degrees Celsius storage; miscibility predicted as favourable (chi equals 0.28, below chi_c equals 0.36); overall confidence 78 percent (high-confidence tier, $P(\text{top-1})$ greater than 0.70); note that this ranking is exploratory at n equals 6 polymers and requires confirmation at n greater than or equal to 20 before confirmatory claims are made.'

The interpretation explicitly states the exploratory status, the recommended next steps (multi-expert AHP confirmation, n expansion to 20, negative-control inclusion in wet-lab), and the conditions under which the framework's recommendation would be revised (if multi-expert W is less than 0.50, if $P(\text{top-1})$ drops below 0.40, or if the held-out test set shows gross ranking failure).

13. Scientific Figures Catalogue (Revised)

13.1 Figure Catalogue (Reduced from 10 to 8 Main-Text Figures)

Version 2.0 reduces the main-text figure count from ten to eight, in response to reviewer criticism 11.1 that the Version 1.0 density (15 figures plus 15 tables in 8,000 to 10,000 words) was too high. Figures 5 (polymer screening funnel) and 10 (CCI component contribution radar) are moved to supplementary material. Two new figures are added: Figure 11 (PCA scree plot and effective dimensionality) and Figure 12 (logistic regression FBM contour with bootstrap CI). The main-text figures are: Figure 1 (Framework Architecture, eight layers), Figure 2 (End-to-End Workflow), Figure 3 (Decision Tree with four gates), Figure 4 (Evidence Integration Flow with PCA), Figure 6 (AHP-TOPSIS Ranking), Figure 7 (Sensitivity Heatmap with Morris), Figure 8 (Uncertainty Propagation), Figure 9 (Validation Workflow with negative controls and baselines), Figure 11 (PCA Scree Plot), Figure 12 (Logistic Regression FBM Contour).

Table 13.1 — Scientific Figures Catalogue (V2.0)

Figure	Title	Status	Purpose
1	Framework Architecture (8 Layers)	Retained	Architecture overview
2	End-to-End Workflow	Retained	11-step workflow
3	Decision Tree	Revised	4 gates (Gate 4 added for FBM)
4	Evidence Integration Flow	Revised	PCA pre-processing added
5	Polymer Screening Funnel	Moved to Supp.	Density reduction
6	AHP-TOPSIS Ranking	Retained	Multi-expert weights
7	Sensitivity Heatmap	Revised	Morris added; Sobol removed
8	Uncertainty Propagation (MC)	Retained	Joint distribution
9	Validation Workflow	Revised	Negative controls + baselines
10	CCI Component Contribution Radar	Moved to Supp.	Redundant with PCA
11	PCA Scree Plot and Effective Dimensionality	NEW	PCA dimensionality reporting
12	Logistic Regression FBM Contour with Bootstrap CI	NEW	FBM formal definition

13.2 Figure Embedding in the Decision Report

The Decision Report embeds Figures 1, 6, 7, 8, 11, and 12 as standard content. Figures 2, 3, 4, and 9 are methodological (documentation only, not in per-decision report). The graphical abstract specification is updated to a minimum of 800 by 800 pixels (criticism 11.8), with 1320 by 1320 pixels preferred for high-resolution journal submission. All figures are generated programmatically by scripts/generate_figures.py and saved as 300 DPI PNG files.

14. Python Implementation Plan (Retained with Module Additions)

14.1 Folder Structure (Retained with New Modules)

The folder structure is retained from Version 1.0: `asd_framework/` (root package), `config/` (`drug_indomethacin.json`, `polymer_library.csv`, `ahp_matrix.json`, `workflow_config.yaml`), `data/` (raw and processed), `src/asd_mcda/` (`knowledge`, `descriptors`, `compatibility`, `integration`, `mcda`, `prediction`, `sensitivity`, `uncertainty`, `validation`, `reporting`, `cli`), `tests/`, `results/`, `notebooks/`, `requirements.txt`, `environment.yml`, `README.md`, `LICENSE (MIT)`. Four new modules are added to `src/asd_mcda/`: `pca.py` (PCA on compatibility signals), `fbm.py` (logistic regression boundary with bootstrap CI), `multi_expert_ahp.py` (geometric-mean aggregation across experts with Kendall W), `baseline_comparison.py` (HSP-only and equal-weight baselines with delta Spearman rho).

14.2 Module Responsibilities and Function Signatures (Revised)

The function signatures from Version 1.0 are retained. New signatures: `compute_pca(scores_matrix, variance_threshold=0.95)` returns (`scores`, `loadings`, `explained_variance`, `k`); `fit_logistic_fbm(X, y, regularization='l2', C=1.0)` returns (`beta`, `intercept`, `auc`); `bootstrap_fbm_ci(X, y, n_boot=10000)` returns (`boundary_ci`); `multi_expert_ahp(matrices)` returns (`consolidated_weights`, `individual_weights`, `CR_individual`, `CR_aggregated`, `kendall_W`); `baseline_comparison(scores, experimental_ranking)` returns (`spearman_hsp_only`, `spearman_equal_weight`, `delta_full_vs_hsp`, `delta_full_vs_equal`).

14.3 Input/Output Schemas (Retained)

The input/output schemas from Version 1.0 are retained: `drug_indomethacin.json`, `polymer_library.csv` (six polymers), `ranking.csv` (output with V2.0 fields: `pca_effective_dimensionality`, `multi_expert_kendall_w`, `baseline_comparison_delta`, `fbm_boundary_logistic`, `fbm_boundary_ci`, `held_out_test_results`, `negative_control_results`), `decision_report.xlsx` (3 sheets: `Summary`, `Ranking`, `Sensitivity`). The `drug_indomethacin.json` now includes `amorphous_density` (measured or estimated) alongside `crystalline_density`, addressing the Simha-Boyer bias (criticism 4.4).

14.4 Reproducibility Plan (Revised with Docker DOI)

The reproducibility plan from Version 1.0 is retained and extended. Pinned dependencies: `rdkit==2024.3.5`, `numpy==1.26.4`, `pandas==2.2.2`, `scipy==1.13.1`, `matplotlib==3.9.0`, `scikit-learn==1.5.1` (new for PCA and logistic regression), `SALib==1.5.1` (new for Morris screening). Random seed: 42 (for Monte Carlo reproducibility). Single-command entry: `python -m asd_mcda.cli --config config/workflow_config.yaml`. The Docker container is deposited on Docker Hub (or GitHub Container Registry) with a Zenodo DOI for archival, addressing reviewer criticism 11.4. The README documents the expected output (paths, sizes, SHA-256 checksums) and the Docker image rebuild frequency (annual, with version-pinned tags).

15. Statistical Framework (Substantially Revised)

15.0 Pre-Registered Statistical Decision Rules (New in V2.0)

The Critical Scientific Review (criticism 3.2) identified that Version 1.0 pre-registered hypotheses on Zenodo but did not pre-register statistical decision rules. Without these, the analysis is post-hoc regardless of when the data are collected. Version 2.0 specifies the seven-element pre-registration for each hypothesis: (a) primary endpoint; (b) test statistic; (c) null and alternative in formal notation; (d) significance level alpha; (e) minimum effect size of interest; (f) stopping rule for data collection; (g) interpretation of borderline results. The full pre-registration template is provided in Appendix D and is deposited on Zenodo before wet-lab work begins (Gate 4 in the workflow).

For H1 (rho greater than or equal to 0.70): primary endpoint is Spearman rho between computational and experimental rankings; test statistic is Spearman rho with 95 percent Fisher-z CI; null is rho less than 0.70; alternative is rho greater than or equal to 0.70; alpha equals 0.05; minimum effect size is rho greater than or equal to 0.70; stopping rule is data collection ceases after pre-specified polymer set is characterised; borderline interpretation is 'exploratory, requires confirmation at n greater than or equal to 20' if CI includes 0.70.

For H2 (RMSE less than or equal to 10 K): primary endpoint is RMSE of Tg prediction; test statistic is RMSE and MAE with 95 percent bootstrap CI; null is RMSE greater than or equal to 10 K; alternative is RMSE less than 10 K; alpha equals 0.05; minimum effect size is RMSE less than 10 K; stopping rule is same; borderline interpretation is 'marginally supported, Gordon-Taylor flagged for Kwei-correction fitting' if RMSE is 9 to 11 K.

For H3 (rank-1 dissolution greater than rank-2): primary endpoint is Q15 and Q30 at pH 6.8; test statistic is two-sample t-test (or Mann-Whitney U if Shapiro-Wilk p less than 0.05) with Cohen's d; null is Q15 and Q30 of rank-1 equals rank-2; alternative is rank-1 greater than rank-2; alpha equals 0.05; minimum effect size is Cohen's d greater than 0.5; stopping rule is same; borderline interpretation is 'inconclusive at n equals 12 vessels, requires n greater than or equal to 24 for confirmation' if p is 0.05 to 0.10.

For H4 (P(failure) less than 0.5 through 6 months): primary endpoint is logistic regression P(failure) at T equals 0, 1, 3, 6 months; test statistic is logistic regression prediction with 95 percent bootstrap CI; null is P(failure) greater than or equal to 0.5 at any timepoint; alternative is P(failure) less than 0.5 at all timepoints; alpha equals 0.05; minimum effect size is P(failure) less than 0.5; stopping rule is same; borderline interpretation is 'Warning region (0.30 to 0.70), boundary flagged for refinement' if any timepoint is 0.30 to 0.70.

15.1 Descriptive Statistics (Retained)

Descriptive statistics (mean, SD, min, max, quartiles) are computed for HSP components, chi, GT-Tg, descriptors, CCI, and TOPSIS CL. These provide the basis for the score normalisation and the candidate-set characterisation.

15.2 Sample Size and Power (Revised)

Version 2.0 acknowledges explicitly that n equals 6 to 8 polymers is underpowered for confirmatory rank-correlation claims. The 95 percent confidence interval for Spearman rho at n equals 6 and rho equals 0.83 spans approximately 0.20 to 0.98 (Fisher z-transformation), which is too wide to confirm or refute H1. The framework treats any H1 result at n equals 6 to 8 as exploratory and specifies n greater than or equal to 20 as the confirmatory threshold (95 percent CI for rho at rho equals 0.83 narrows to approximately 0.60 to 0.93 at n equals 20). The computational expansion to n equals 20 polymers (literature-only, no wet-lab) is specified in the implementation roadmap (Section 23, Phase 2) as a mandatory step before any confirmatory claim.

15.3 Multiple Comparisons (Revised with Family Specification)

The multiple-comparison corrections from Version 1.0 are retained, with explicit specification of the multiplicity family (criticism 6.4). Tukey HSD is applied to the family of all pairwise polymer comparisons within a single analysis (e.g., dissolution Q30 across polymers at a single pH and timepoint; family size equals $k(k-1)/2$ where k is the number of polymers). Dunnett is applied to the family of comparisons of each polymer against a reference (e.g., each polymer vs Soluplus; family size equals $k-1$). Benjamini-Hochberg FDR (q equals 0.05) is applied to the family of descriptor screening comparisons (family size equals number of descriptors tested, typically 8). Each family is specified in the pre-registration; no cross-family correction is applied.

15.4 Regression and Residual Analysis (Retained)

Linear regression of predicted Tg on experimental Tg with identity y equals x reference is retained. Residual plot and normal QQ plot are reported. If non-normal (Shapiro-Wilk p less than 0.05), bootstrap CI is reported instead of parametric CI.

15.5 Confidence Intervals (Retained with Permutation Test Addition)

95 percent confidence intervals for predicted Tg (from Monte Carlo) and for validation metrics (RMSE, MAE, Spearman rho) are computed via non-parametric bootstrap with 10,000 resamples. Bootstrap is preferred because the small sample size (typically 6 drug-polymer pairs) makes parametric assumptions unreliable. For paired dissolution profiles (f2 calculation), the permutation test is added as an alternative to bootstrap (criticism 6.3): the permutation test preserves the pairing by permuting test/reference labels within each vessel, producing a tighter CI than the bootstrap. Both bootstrap and permutation CIs are reported; the permutation CI is preferred for paired-profile analyses.

15.6 R-squared Predicted Threshold (Revised from 0.50 to 0.70)

The R-squared predicted (R^2_{pred}) acceptance criterion is revised from 0.50 (Version 1.0) to 0.70 (criticism 6.5). R^2_{pred} equals 0.50 means the model explains 50 percent of variation in new observations, which is modest predictive power. For Q1 publication, R^2_{pred} greater than or equal to 0.70 is more appropriate. The R^2_{pred} must also be within 0.20 of the adjusted R^2 (which must be greater than or equal to 0.70). Q-squared (cross-validated R^2) is now reported alongside R^2_{pred} (criticism Q-MOD-024).

15.7 Bayesian Sensitivity Analysis (New in V2.0)

The Critical Scientific Review (criticism 6.7) identified that the blanket exclusion of Bayesian methods in Version 1.0 was unjustified for small-sample stability data. Version 2.0 adds a Bayesian sensitivity analysis parallel to the frequentist primary analysis. For stability data (n equals 3 timepoints, single batch), a weakly-informative prior (normal(0, 10) on the intercept, normal(0, 5) on the slope) is placed on the linear trend regression of assay or Tg against time. The posterior is computed via Markov Chain Monte Carlo (PyMC or Stan) with 4 chains, 2000 iterations each, and $\hat{R}$ less than 1.01. The posterior credible interval (95 percent) for the trend slope is reported alongside the frequentist CI. If the Bayesian and frequentist CIs agree, the result is reported as 'robust'; if they disagree, the result is reported as 'sensitive to prior choice' and the discrepancy is investigated. ICH Q1E does not preclude Bayesian methods; the FDA has accepted Bayesian stability analyses since 2010.

15.8 Equivalence Testing for FBM (New in V2.0)

The Two One-Sided Tests (TOST) procedure for the FBM boundary is specified in Section 11.5. The pre-specified margin is plus or minus 5 percent dissolution. The null hypothesis is that Q30 is outside the equivalence interval (less than 75 percent or greater than 85 percent); the alternative is that Q30 is within the interval. If the null is rejected (both one-sided p-values less than 0.05), the boundary is confirmed at that operating point.

15.9 Pseudoreplication Acknowledgment (Revised)

The Critical Scientific Review (criticism 5.3, 6.9) identified that a mixed-effects model with a single batch is mathematically degenerate (the random effect has zero degrees of freedom) and is therefore an admission, not a mitigation. Version 2.0 specifies three independent spray-drying batches for the rank-1 polymer (mandatory change 3, Section 16), allowing a genuine random-effect estimate. If three batches are logistically infeasible, the framework falls back to a single-batch analysis with an explicit sensitivity analysis: the single-batch result is compared against the literature-reported batch-to-batch variability for the same polymer-drug combination, and the conclusion is reported as 'single-batch; sensitivity analysis suggests the conclusion is (robust / sensitive) to batch effects.' This honest acknowledgment replaces the Version 1.0 claim that a mixed-effects model 'mitigates' single-batch pseudoreplication.

16. Experimental Framework (New in V2.0)

16.1 Characterization Methods (Revised)

DSC (Differential Scanning Calorimetry) is used for T_g measurement and single-T_g miscibility assessment. Standard conditions: 10 degrees Celsius per minute heating rate, nitrogen purge 50 mL per minute, hermetically sealed aluminium pans, sample mass 3 to 5 mg. Modulated DSC (MDSC) is recommended (criticism Q-MOD-011) to separate reversing and non-reversing heat flow, particularly for samples with residual solvent evaporation or relaxation endotherms. MDSC conditions: underlying heating rate 2 degrees Celsius per minute, modulation amplitude plus or minus 0.5 degrees Celsius, modulation period 60 seconds.

FTIR (Fourier Transform Infrared Spectroscopy) is used for hydrogen-bonding assessment via carbonyl stretch shifts. Indomethacin C=O stretch at 1715 cm⁻¹ (carboxylic acid) and 1690 cm⁻¹ (amide); polymer C=O varies (PVP K30 at 1660 cm⁻¹, PVP-VA 64 at 1680 cm⁻¹, HPMCAS-L at 1740 cm⁻¹). ATR-FTIR is preferred over KBr pellet (criticism Q-MOD-012) for greater sensitivity to surface composition and avoidance of KBr hydration artefacts. Spectral resolution 4 cm⁻¹, 64 scans co-added.

PXRD (Powder X-Ray Diffraction) is mandatory for ALL top-2 ASDs at T equals 0 minimum (criticism 5.2, mandatory change 4), with T equals 6 months added for any ASD selected for stability. PXRD conditions: Cu-Kalpha radiation, 40 kV, 40 mA, 2-theta range 5 to 40 degrees, step size 0.02 degrees, scan speed 2 degrees per minute. The absence of crystalline peaks confirms amorphicity; the presence of crystalline peaks indicates phase separation or recrystallisation. PXRD cannot detect amorphous-amorphous phase separation (AAPS) at domain sizes below 20 nm; this limitation is acknowledged and solid-state NMR is noted as the gold standard but out of scope for M.Pharm.

16.2 Spray Drying (Revised with 3-Batch Requirement)

Spray drying is performed on a Buchi B-290 lab-scale spray dryer with inert loop for solvent recovery. Standard conditions: inlet temperature 80 to 120 degrees Celsius (DoE range), feed rate 5 to 15 mL per minute (DoE range), atomisation nitrogen flow 600 L per hour, aspirator rate 100 percent. Solvent: acetone (preferred for indomethacin solubility and low boiling point) or ethanol (alternative; Q-MOD-010 to be resolved in favour of acetone based on indomethacin solubility). Drug loading: 30 percent w/w anchor (justified as typical ASD range for indomethacin, criticism Q-MOD-009), with the loading-range analysis (10 to 50 percent w/w in 5 percent steps) for the rank-1 polymer.

Three independent spray-drying batches are produced for the rank-1 polymer (mandatory change 3, criticism 5.3, 6.9) to enable genuine batch-to-batch variance estimation. Each batch is produced on a different day with fresh feed solution. The three batches are characterised by DSC, FTIR, PXRD, and dissolution. If three batches are logistically infeasible within M.Pharm scope, the framework falls back to a single-batch analysis with an explicit sensitivity analysis (Section 15.9).

The Box-Behnken DoE is retained (17 runs per polymer) but with 5 centre points (revised from 3, criticism Q-MOD-003) to provide 6 pure-error degrees of freedom (revised from 4) for adequate lack-of-fit power. The DoE factors are: drug-polymer ratio (20 to 40 percent w/w), inlet temperature (80 to 120 degrees Celsius), feed concentration (5 to 15 percent w/v).

16.3 Tablet Validation (Revised with 5 Centre Points)

Tablet compression is performed on a single-punch tablet press (Korsch XP-1 or equivalent) with 8 mm round flat-faced tooling. The 2-cubed factorial DoE is retained (8 factorial runs) but with 5 centre points (revised from 3, criticism Q-MOD-003) and 4 axial points (Central Composite Design, criticism Q-MOD-004) for better curvature detection. The DoE factors are: compression force (10 to 20 kN), filler fraction (30 to 50 percent w/w), and disintegrant fraction (2 to 6 percent w/w). Tablet CQAs: hardness (4 to 8 kp, USP <1216>), friability (less than 1 percent, USP <1216>), disintegration (less than 15 minutes, USP for IR), content uniformity (95 to 105 percent per USP <905>; ICH Q4B alignment noted, criticism Q-MOD-019), assay (95 to 105 percent).

Tablet formulation: SD-ASD (30 percent drug load equivalent), microcrystalline cellulose (MCC) filler, croscarmellose sodium disintegrant, magnesium stearate lubricant. Excipient compatibility with SD-ASD is verified by DSC and FTIR (criticism Q-MOD-015). Tablet-to-tablet weight variation is added as an additional CQA (criticism Q-MOD-017).

16.4 Dissolution (Revised: Multi-pH Mandatory)

Dissolution testing is performed per USP <711> Apparatus II (paddles), 900 mL medium, 50 rpm, 37 plus or minus 0.5 degrees Celsius. Multi-pH media are MANDATORY (criticism 5.4, optional improvement 3): pH 1.2 (0.1 N HCl, gastric), pH 4.5 (acetate buffer, intermediate), and pH 6.8 (phosphate buffer, intestinal). Sampling timepoints: 1, 2, 5, 10, 15, 30, 45, 60 minutes (early timepoints added, criticism Q-MOD-013). Twelve vessels per polymer per pH medium (retained from V1.0; the reviewer suggested 6 USP minimum or 24 for enhanced power; 12 is retained as a balance).

Dissolution profile comparison: f2 similarity factor (Shah and co-workers 1998) for test vs reference (rank-1 vs rank-2, or ASD vs crystalline). Bootstrap CI (n_boot equals 10,000) and permutation test CI (Section 15.5) are reported. Spring-and-parachute classification: Type I (spring-and-parachute), Type II (spring-only), Type III (no spring) based on Tg margin and chi.

16.5 Stability (Revised: 6-Month Accelerated + 12-Month Future)

Stability testing is performed per ICH Q1A(R2). The 6-month accelerated stability (40 degrees Celsius plus or minus 2, 75 percent RH plus or minus 5) is retained as the M.Pharm-feasible study. Stability CQAs: assay (HPLC-UV), dissolution (multi-pH Q15, Q30, Q60), Tg (DSC), FTIR (carbonyl shift), PXRD (crystallinity), water content (Karl Fischer, added per criticism Q-MOD-020), residual solvent (GC-HS, added per criticism 5.6). Timepoints: 0, 1, 3, 6 months.

The 12-month long-term stability (25 degrees Celsius plus or minus 2, 60 percent RH plus or minus 5) and the 6-month intermediate stability (30 degrees Celsius plus or minus 2, 65 percent RH plus or minus 5) are specified as PLANNED FUTURE WORK (criticism 5.5, mandatory change 5). The committee acknowledges that 6-month accelerated data alone are the ICH Q1A minimum for NDA submission but are insufficient for Q1 publication; 12-month long-term data are required for Q1. The 12-month study is initiated in parallel with the M.Pharm programme but is completed after M.Pharm submission; the data are reported in a follow-up publication or in the M.Pharm thesis appendix if available.

ICH Q1E statistical evaluation procedure (criticism 10.7): the trend is evaluated at the batch level (not pooled across batches, given the small number of batches), using linear regression of assay or Tg against time. The extrapolation model is linear (with Arrhenius correction for temperature). The confidence level is 95 percent one-sided. The shelf-life is the time at which the 95 percent lower confidence bound on assay crosses 95 percent (the lower acceptance criterion).

16.6 Analytical Methods (Revised: GC-HS, HPLC-UV Q2(R2), SEM/PSA Added)

Three analytical methods are added to address reviewer criticism 5.6 and regulatory criticism 10.4. First, residual solvent testing by GC-HS (Gas Chromatography with Headspace Sampler) per ICH Q3C: column DB-624 (30 m, 0.32 mm, 1.8 micrometre), carrier gas helium, oven programme 40 degrees Celsius (5 min) then 10 degrees Celsius per minute to 240 degrees Celsius (5 min), FID detector 250 degrees Celsius, headspace incubation 80 degrees Celsius for 30 minutes. Limit: Class 3 solvents (acetone PDE 50 mg/day; ethanol PDE 50 mg/day).

Second, HPLC-UV for indomethacin assay per ICH Q2(R2): column C18 (250 x 4.6 mm, 5 micrometre), mobile phase acetonitrile-phosphate buffer pH 3.5 (60:40), flow 1.0 mL/min, detection 254 nm, injection 20 microlitre, runtime 10 min. Q2(R2) validation parameters: linearity (5 concentrations, 50 to 150 percent of nominal), accuracy (recovery 98 to 102 percent), precision (repeatability RSD less than 2 percent; intermediate precision RSD less than 3 percent), LOD (S/N greater than 3), LOQ (S/N greater than 10), specificity (no interference from polymer or excipients), robustness (deliberate variation of flow, mobile phase composition, column temperature).

Third, particle size and morphology by SEM (Scanning Electron Microscopy) and laser diffraction (PSA): SEM (JEOL JSM-6390 or equivalent) at 5 kV, gold sputter coating 60 seconds, magnification 500x to 5000x. Laser diffraction (Malvern Mastersizer 3000 or equivalent) with dry dispersion, refractive index 1.52 for indomethacin, volume-weighted D10, D50, D90 reported. These methods are RECOMMENDED (optional improvement 5) but add substantial publication value.

ICH Q2(R2) validation is extended to DSC (precision: RSD less than 0.5 degrees Celsius for Tg of reference standard), FTIR (precision: peak position reproducibility less than 1 cm⁻¹), and dissolution (precision: RSD less than 5 percent at Q15, less than 3 percent at Q60), addressing criticism 10.4.

17. Worked Example: Indomethacin + 6 Polymers (Revised)

17.1 Drug Knowledge (Indomethacin)

Indomethacin property sheet (retained from V1.0, with amorphous density added per criticism 4.4). IUPAC: 2-{1-[(4-chlorophenyl)carbonyl]-5-methoxy-2-methyl-1H-indol-3-yl}acetic acid. Formula C₁₉H₁₆ClNO₄. MW 357.79 g/mol. Canonical SMILES: CC1=C(C=C(C=C1)OC)C2=C(C3=CC=CC=C3N2CC(=O)O)C(=O)C4=CC=C(C=C4)Cl. CAS 53-86-1. pKa 4.5 (carboxylic acid). logP 4.27 (RDKit Crippen). logD (pH 7.4) 1.32. Tm (gamma-form) 151 degrees Celsius equals 424.15 K (Yamamoto 1996). Tm (alpha-form) 154 degrees Celsius equals 427.15 K (Chavan 2017). Tg (amorphous) 42 degrees Celsius equals 315.15 K (Baird 2010). Tg (Boyer-Beaman estimate) 297 K (0.7 times 424). Density (crystalline) 1.31 g/cm³. Density (amorphous, NEW in V2.0) 1.22 g/cm³ (estimated; measured value to be confirmed experimentally per criticism Q-MAJ-039). Delta-H-fus 34.5 kJ/mol. Aqueous solubility 5 micrograms per mL at pH 1.2; 100 micrograms per mL at pH 6.8. BCS Class II. HBD 2, HBA 4. TPSA 68.5 square angstroms. Aromatic rings 2. Rotatable bonds 4. HSP delta_D 19.2, delta_P 7.9, delta_H 8.4 MPa^{0.5} (Hansen 2007). R_o 8.0 MPa^{0.5}. Molar volume V_m 273 cm³/mol.

17.2 Polymer Library (Retained)

The six-polymer library is retained: PVP K30 (neutral, N-vinylpyrrolidone, Mn 40 kDa, Tg 443 K, density 1.20, delta_D 17.4, delta_P 8.2, delta_H 11.7); PVP-VA 64 (neutral copolymer VP/VA 6:4, Mn 45 kDa, Tg 380 K, density 1.20, delta_D 17.0, delta_P 8.0, delta_H 10.0); HPMCAS-L (anionic cellulosic, Mn 100 kDa, Tg 394 K, density 1.28, delta_D 18.0, delta_P 8.2, delta_H 10.5); Eudragit L100 (anionic methacrylate, Mn 125 kDa, Tg 438 K, density 1.25, delta_D 16.5, delta_P 7.5, delta_H 9.0); Soluplus (amphiphilic graft PCL-PVAc-PEG, Mn 90 kDa, Tg 343 K, density 1.15, delta_D 18.0, delta_P 8.5, delta_H 10.5); HPMC E5 (neutral cellulosic, Mn 20 kDa, Tg 438 K, density 1.33, delta_D 18.5, delta_P 8.8, delta_H 12.0).

Table 17.1 — Polymer Library (6 Polymers, Retained from V1.0)

Polymer	Class	Mn (kDa)	Tg (K)	rho (g/cm ³)	delta_D	delta_P	delta_H	Lit. w/ INM
PVP K30	Neutral	40	443	1.20	17.4	8.2	11.7	Miscible (Taylor 2007)
PVP-VA 64	Neutral	45	380	1.20	17.0	8.0	10.0	Miscible (Van Eerden brugh 2010)
HPMCAS-L	Anionic	100	394	1.28	18.0	8.2	10.5	Miscible (Ueda 2014)
Eudragit L100	Anionic	125	438	1.25	16.5	7.5	9.0	Partially (Rumondor 2010)
Soluplus	Amphiphilic	90	343	1.15	18.0	8.5	10.5	Miscible (Hardung 2010)
HPMC E5	Neutral	20	438	1.33	18.5	8.8	12.0	Miscible (Tres 2015)

17.3 Compatibility Computation (Retained)

The compatibility computation is retained from Version 1.0. At 30 percent w/w drug loading, the six polymers produce the following scores: Soluplus has the lowest R_a (3.2), RED (0.40), and χ (0.28); HPMCAS-L is second lowest; Eudragit L100 has the highest R_a (5.4) and χ (0.51). All six pass the HSP gate (RED less than 1.0) and the χ gate (χ less than or equal to 0.5). The Gordon-Taylor T_g predictions range from 338 K (Soluplus, 65 degrees Celsius) to 410 K (PVP K30, 137 degrees Celsius). The Lindvig parameters V_{m_drug} equals 273 cm³/mol and T equals 298 K are used.

Table 17.2 — Compatibility Computation at 30% w/w Drug Loading (Retained)

Polymer	R_a	RED	χ	K (SB)	T_{g_mix} (K)	s_{HSP}	s_{χ}	s_{GT}
PVP K30	4.6	0.58	0.42	0.99	410	0.71	0.58	0.94
PVP-VA 64	3.8	0.48	0.34	1.18	367	0.76	0.66	0.66
HPMCAS-L	3.4	0.43	0.30	1.04	372	0.79	0.70	0.71
Eudragit L100	5.4	0.68	0.51	1.00	406	0.66	0.49	0.92
Soluplus	3.2	0.40	0.28	1.34	338	0.80	0.72	0.42
HPMC E5	4.5	0.56	0.40	1.16	402	0.72	0.60	0.91

17.4 AHP Pairwise Matrix and Weights (Revised: Multi-Expert)

Version 2.0 specifies multi-expert AHP elicitation as the primary method. For the worked example, a panel of three pharmaceutical scientists (formulation scientist, analytical chemist, regulatory scientist) independently completes the pairwise comparison matrix on the principal components (after PCA, see Section 17.5). The default matrix from Version 1.0 is retained as the starting point for individual elicitation. The geometric-mean aggregation produces the consolidated weight vector. Inter-expert agreement (Kendall W) and individual CR values are reported.

The default matrix from Version 1.0 (with weights w_{HSP} 0.28, w_{χ} 0.30, w_{desc} 0.12, w_{GT} 0.18, w_{lit} 0.12; CR 0.07) is retained as the fallback single-expert matrix. In the multi-expert setting, this matrix is the starting point; individual experts revise the pairwise judgements based on their domain expertise, and the consolidated weights are computed. The consolidated weights may differ from the V1.0 defaults; the worked example in V2.0 reports the consolidated weights as 'to be determined by multi-expert elicitation (Section 23, Phase 2).'

17.5 PCA and CCI Computation (Revised)

PCA on the five-score matrix (6 polymers by 5 scores) is performed after centering and scaling. Given the high HSP- χ correlation (ρ equals 0.94 in V1.0), PC1 is expected to capture the cohesive-energy cluster (HSP, χ , GT loadings) and to explain a large fraction of variance. The expected output is: PC1 explains approximately 65 to 75 percent of variance (dominated by HSP, χ , GT loadings); PC2 explains approximately 15 to 20 percent (dominated by descriptors and literature loadings); PC3 through PC5 explain the remaining 5 to 15 percent. The number of retained PCs (k) for the 95 percent variance threshold is expected to be 2.

If PC1 explains more than 80 percent of variance (the threshold for 'effectively one-dimensional' classification), the CCI is reported as 'effectively one-dimensional, dominated by cohesive-energy compatibility.' If PC1 explains 65 to 80 percent, the CCI is 'two-dimensional, with cohesive-energy and structure-based components.' The actual PCA output will be computed when the framework is implemented; the worked example in V2.0 reports the expected PCA structure qualitatively and notes that the quantitative output will be generated by the Python implementation.

The CCI is computed as the weighted combination of the retained PCs, with weights from multi-expert AHP applied to the PCs. The ranking may shift slightly from the Version 1.0 ranking (Soluplus rank-1, PVP-VA 64 rank-2, PVP K30 rank-3) due to the PCA pre-processing. The V2.0 worked example reports the expected ranking as 'to be confirmed by implementation' but retains the V1.0 ranking as the provisional reference.

17.6 Baseline Comparison (New in V2.0)

The baseline comparison (Section 11.4) is applied to the worked example. HSP-only ranking (ranking by s_{HSP} alone): Soluplus (0.80), HPMCAS-L (0.79), PVP-VA 64 (0.76), HPMC E5 (0.72), PVP K30 (0.71), Eudragit L100 (0.66). Equal-weight ranking (ranking by simple average of all five scores, weights 0.20 each): expected to be similar to the full CCI ranking but with smaller deltas between polymers. The full CCI ranking (with PCA and AHP) is the V2.0 ranking. The delta Spearman rho between the full CCI and the HSP-only baseline, and between the full CCI and the equal-weight baseline, will be computed when the framework is implemented. If the full CCI does not outperform the baselines by delta rho greater than or equal to 0.10, the committee recommends simplifying the framework (Section 11.4).

17.7 LOO Cross-Validation (New in V2.0)

Leave-one-out cross-validation (LOO CV) is performed on the six-polymer set. Each polymer is held out in turn; the AHP weights are derived from the remaining five polymers; the held-out polymer is ranked; the rank is compared against the rank obtained when all six polymers are used. LOO CV provides a more honest estimate of generalisation error than in-sample validation. The expected LOO CV result is that the ranking is stable (held-out polymer is assigned the same rank as in the full-set analysis) for most polymers, with possible rank swaps between adjacent polymers (e.g., PVP-VA 64 and PVP K30 swapping rank-2 and rank-3). The LOO CV results will be computed when the framework is implemented.

17.8 Held-Out Test Set (New in V2.0)

The held-out test set protocol (Section 11.2) is applied: AHP weights are derived from four polymers (PVP K30, PVP-VA 64, HPMCAS-L, Soluplus) and the framework is tested on two unseen polymers (Eudragit L100, HPMC E5). The expected result is that Eudragit L100 is correctly ranked low (rank-5 or 6) and HPMC E5 is correctly ranked in the lower half (rank-4 to 6). If the held-out polymers are grossly mis-ranked (e.g., Eudragit L100 ranked rank-1), this indicates a failure of the framework's specificity and triggers a refinement. The held-out test set results will be computed when the framework is implemented.

17.9 Validation Against Literature (Revised)

The literature validation is retained from Version 1.0 with explicit exploratory classification. Taylor and co-workers (2007) and Van Eerdenbrugh and co-workers (2010): indomethacin-PVP K30 and indomethacin-PVP-VA 64 miscible (single-Tg by DSC), consistent with framework top-3. Hardung (2010) and Nagy and co-workers (2012): indomethacin-Soluplus miscible with significant Tg elevation, consistent with rank-1. Rumondor and co-workers (2010): indomethacin-Eudragit L100 partially miscible with tendency to phase separation, consistent with rank-6. The qualitative agreement between the framework's ranking and the published experimental ranking is Spearman rho equals 0.83.

Version 2.0 explicitly states that the 95 percent confidence interval for rho at n equals 6 and rho equals 0.83 spans approximately 0.20 to 0.98 (Fisher z-transformation), which is too wide to confirm or refute H1. The rho equals 0.83 result is therefore marked as exploratory, requiring confirmation at n greater than or equal to 20. The computational expansion to n equals 20 polymers (literature-only, no wet-lab) is specified in the implementation roadmap (Section 23, Phase 2). This honest acknowledgment replaces the Version 1.0 claim

that rho equals 0.83 is 'excellent agreement.'

18. Failure Boundary Mapping (New in V2.0)

18.1 Formal Definition (Logistic Regression)

The Failure Boundary Map (FBM) is formally defined via logistic regression. The binary outcome variable is 'operational failure' (defined as dissolution Q30 less than 80 percent at pH 6.8, OR Tg shift exceeding 5 K indicating phase separation, OR appearance of crystalline peaks by PXRD). The predictors are the four DoE factors: polymer rank (1, 2, 3, ...), inlet temperature (degrees Celsius), drug loading (percent w/w), and feed concentration (percent w/v). The model is: $\text{logit}(P(\text{failure})) = \beta_0 + \beta_1 \text{polymer_rank} + \beta_2 \text{inlet_T} + \beta_3 \text{drug_loading} + \beta_4 \text{feed_conc}$. The failure boundary is the contour where $P(\text{failure}) = 0.5$.

This formal definition replaces the Version 1.0 operational threshold (Q30 less than 80 percent) and addresses reviewer criticisms 7.5, 8.1, and 8.5. The logistic regression provides a continuous probability of failure rather than a binary classification, enabling a probabilistic Safe/Warning/Failure taxonomy. The terminology uses 'operational failure' or 'CQA failure' rather than 'failure' alone, addressing criticism 8.5 that the latter overstates the consequences.

18.2 Safe/Warning/Failure Taxonomy (Quantified)

The Safe, Warning, and Failure regions are defined probabilistically: Safe is $P(\text{failure}) < 0.30$, Warning is $0.30 \leq P(\text{failure}) < 0.70$, and Failure is $P(\text{failure}) \geq 0.70$. This quantifies the Warning region (which was undefined in Version 1.0's 'approaches limit' language) and provides a continuous gradient between safe and failing operating conditions. The 0.30 and 0.70 thresholds are operationally justified: 0.30 corresponds to a less than 1-in-3 chance of failure (acceptable for a pre-screen), and 0.70 corresponds to a greater than 2-in-3 chance of failure (unacceptable for wet-lab commitment).

18.3 Multidimensional Boundary (4D with 2D Slices)

The FBM is a four-dimensional surface (polymer rank, inlet T, drug loading, feed concentration) and cannot be directly visualised. Version 2.0 specifies 2D slices at multiple combinations of the held-fixed factors. For example, the polymer rank versus inlet T slice is plotted at drug loading equals 30 percent w/w and feed concentration equals 10 percent w/v (the DoE centre). Additional slices are plotted at drug loading equals 20 percent and 40 percent, and at feed concentration equals 5 percent and 15 percent, providing a total of 9 slices that characterise the 4D surface. Each slice is a contour plot of $P(\text{failure})$ with the 0.5 contour highlighted as the boundary. This addresses reviewer criticism 8.2 that the Version 1.0 2D contour ignored two DoE factors.

18.4 Boundary Estimation (Minimum 20-30 Points)

The minimum number of data points for boundary estimation is 20 to 30 (criticism 8.3). The Version 1.0 approximately 5 boundary points is statistically indefensible. The Box-Behnken DoE (17 runs per polymer) plus 5 centre points provides 22 runs per polymer; for two polymers (rank-1 and rank-2), this is 44 runs, which is sufficient for logistic regression with 4 predictors (10:1 ratio of observations to predictors). The boundary is estimated using maximum likelihood (scikit-learn LogisticRegression with L2 regularisation, C equals 1.0). Model fit is assessed by AUC-ROC; an AUC greater than or equal to 0.75 is required for the boundary to be reported as actionable.

18.5 Pre-Registration of Boundary Location

The boundary location is pre-registered on Zenodo before wet-lab validation begins (Gate 4 in the workflow). The pre-registration includes: (a) the logistic regression model specification; (b) the DoE factors and ranges; (c) the binary outcome definition; (d) the planned boundary estimation method (maximum likelihood); (e) the planned bootstrap CI computation (n_{boot} equals 10,000). This pre-registration addresses reviewer criticism 8.6 that post-hoc boundary definition is cherry-picking. The pre-registered boundary is the computational prediction; the wet-lab data are then used to confirm or refute the pre-registered boundary.

18.6 Uncertainty Quantification on Boundary Location

The uncertainty in the boundary location is quantified via bootstrap (Equation 12). The logistic regression is refit on B equals 10,000 bootstrap resamples of the original data, producing a distribution of boundary locations. The 95 percent confidence interval on the boundary is the 2.5th and 97.5th percentiles of the bootstrap distribution. If the CI width exceeds 20 percent of the design-space range in either dimension, the boundary is flagged as 'low confidence' and additional data points are requested before the boundary is used for go/no-go decisions. A Bayesian alternative (placing weakly-informative priors on the beta coefficients and computing the posterior on the boundary location via MCMC) is specified as a recommended extension.

18.7 Terminology: Operational Failure

The term 'failure' is replaced throughout by 'operational failure' or 'CQA failure' (criticism 8.5). A product with Q_{30} equals 79 percent is not meaningfully different from one with Q_{30} equals 81 percent, but the Version 1.0 FBM classified them as failure and success respectively. The Version 2.0 probabilistic FBM assigns them $P(\text{failure})$ values around 0.5, placing both in the Warning region. This honest terminology avoids overstating the consequences of operational failure and aligns with the ICH Q8 design-space concept (where the edge of the design space is a regulatory boundary, not a safety boundary).

18.8 Comparison to ICH Q8 Design Space

The FBM is conceptually similar to the ICH Q8 design-space edge but uses different terminology. The ICH Q8 design space is a multidimensional input region where product quality is assured; the edge is the boundary of this region. The FBM is the contour where $P(\text{failure})$ equals 0.5, which is conceptually equivalent to the design-space edge if 'quality is assured' is interpreted as ' $P(\text{failure})$ less than 0.5.' Version 2.0 explicitly compares the FBM to the design-space edge and acknowledges that the FBM provides no additional information beyond the design-space edge; the contribution is the operationalisation via logistic regression, not a new concept. The reviewer's criticism 8.4 is accepted.

19. Reviewer Simulation (Revised)

19.1 Reviewer Report 1: AAPS PharmSciTech (Primary Target, NEW)

AAPS PharmSciTech is the primary submission target (Section 21). The simulated reviewer is a QbD-focused formulation scientist with 15 years of industrial experience. Major critiques anticipated: (1) the 6-month accelerated stability is the ICH Q1A minimum; 12-month long-term is required for Q1 publication — response: 12-month is initiated in parallel, will be reported in follow-up; (2) the single-batch pseudoreplication is acknowledged — response: three batches for rank-1 polymer (mandatory change 3); (3) the PXRD coverage for top-2 is essential — response: PXRD for all top-2 at T equals 0 (mandatory change 4). Minor comments: add spray-drying process flow diagram; clarify chi units (dimensionless) and temperature dependence (298 K); add glossary.

19.2 Reviewer Report 2: International Journal of Pharmaceutics (Secondary Target)

IJP is the secondary submission target. The simulated reviewer is a senior formulation scientist with 20 years of ASD experience. Major critiques anticipated: (1) the validation scope is retrospective (literature-based) and prospective validation is needed — response: prospective wet-lab validation is the core of the 40-week roadmap; (2) the chi estimation error may be larger for drug-polymer systems — response: chi uncertainty increased to plus or minus 25 percent, MPD calibration recommended; (3) the AHP subjectivity is a concern — response: multi-expert AHP is now primary (mandatory change 7). Minor comments: add spray-drying flow diagram; clarify chi units; add glossary.

19.3 Reviewer Report 3: Molecular Pharmaceutics (Tertiary, Scope-Mismatch Acknowledged)

Molecular Pharmaceutics is the tertiary target, with explicit acknowledgment that the scope is mismatched (the journal favours molecular-level mechanistic studies with strong physicochemical characterisation; the framework's QbD-FBM focus is closer to AAPS PharmSciTech or IJP scope). If submitted to Molecular Pharmaceutics, the anticipated critiques are: (1) lack of MD simulation to validate chi — response: MD is out of M.Pharm scope, future work; (2) HSP-chi cross-correlation double-counts — response: PCA pre-processing (mandatory); (3) the descriptor score is opaque — response: sub-weights specified (HBD 0.3, HBA 0.3, TPSA 0.2, aromatic ratio 0.2). The committee recommends against submitting to Molecular Pharmaceutics unless the computational chemistry depth is substantially increased (e.g., MD validation of chi, which is out of scope).

19.4 Synthesis of Reviewer Feedback (Revised)

The four recurring themes from Version 1.0 are addressed: (i) validation scope limited and should be prospective — addressed by the 40-week roadmap; (ii) HSP-chi cross-correlation should be addressed — addressed by PCA pre-processing; (iii) framework should extend beyond single drug loading — addressed by the loading-range analysis (10 to 50 percent w/w); (iv) bioperformance prediction is a major gap — acknowledged as future work (would require Noyes-Whitney plus PBPK, out of scope). The Version 2.0 framework is considered publishable at AAPS PharmSciTech (primary) or IJP (secondary) after the 40-week roadmap is completed. Molecular Pharmaceutics is not recommended unless scope is increased.

19.5 External Pre-Submission Review (New in V2.0)

The Critical Scientific Review (criticism 11.9) identified that the Version 1.0 reviewer simulation was internal (same committee that wrote the framework). Version 2.0 specifies that at least one external expert review is obtained before submission. The external expert is a Q1-published pharmaceutical scientist not affiliated with the project, identified through the supervisor's professional network or through the journal's reviewer database (if available). The external review is solicited after the 40-week roadmap is completed and the manuscript is drafted; the external reviewer's comments are addressed in a revision before submission. The external review is documented in the supplementary material.

20. Industrial Perspective (Revised)

20.1 Academic Contribution (Current Scope)

The framework's academic contribution is the methodological protocol for rational polymer selection in SD-ASDs, demonstrated on indomethacin. The contribution includes: (a) the Composite Compatibility Index as a transparent integration device with PCA pre-processing; (b) the operationalisation of failure boundary mapping for indomethacin SD-ASDs via logistic regression; (c) the open-source Python and R release with Docker containerisation; (d) the multi-expert AHP elicitation protocol; (e) the pre-registered statistical decision rules. These contributions are academically defensible and publishable in AAPS PharmSciTech or IJP after empirical validation.

20.2 Industrial Applicability (Limited)

The framework's industrial applicability is LIMITED. The framework operates at TRL 4 (laboratory validation) and IRL 2 to 3 (validated in a laboratory setting but not in an industrial pipeline). Industrial adoption would require: (a) bridging to TRL 6 via pilot-scale spray drying (ProCepT or GEA Pharma-Verifier); (b) GMP batch documentation and deviation logs; (c) computer system validation (CSV) for the computational pipeline, taking 6 to 12 months and 200,000 to 500,000 USD; (d) a cost-of-goods (CoG) model; (e) IP freedom-to-operate (FTO) analysis; (f) comparison to COSMO-RS, GastroPlus, Simcyp; (g) a regulatory pathway for the FBM (recognition by regulators as a valid design-space representation). None of these are achievable within M.Pharm scope.

The committee endorses the reviewer's assessment (criticism 9.1 through 9.9) that the framework is not industrial-ready. An industrial formulation scientist would use the framework as a Tier-1 screening tool (pre-IND stage) where the cost of wet-lab screening is high relative to the project budget, but would not use it as a Tier-3 (late-development) tool. The framework's industrial verdict is therefore: academically defensible, industrially limited, future industrial validation required.

20.3 Future Industrial Validation (Deferred)

Future industrial validation is deferred to PhD-scale or industrial-collaboration work. The specific items are: (1) bridge to TRL 6 via pilot-scale spray drying; (2) GMP batch documentation and CSV; (3) CoG model; (4) FTO analysis; (5) comparison to COSMO-RS, GastroPlus, Simcyp; (6) regulatory pathway for the FBM; (7) scale-up bridge from Buchi B-290 to pilot/production scale; (8) multi-drug generalisation across 5 to 10 BCS Class II drugs. These items are out of M.Pharm scope and are not claimed in Version 2.0.

20.4 Industrial Verdict (Revised)

The industrial verdict is: TRL 4 (laboratory validation); IRL 2 to 3 (validated in laboratory, not in industrial pipeline); industrial use as Tier-1 pre-IND screening tool only; not Tier-3 late-development tool. This verdict is honestly stated and replaces the Version 1.0 language that presented the framework as a candidate for industrial adoption. The committee acknowledges that the framework's open-source release creates IP disclosure risk for industrial sponsors (criticism 9.8) and recommends dual-licensing (open-source for academic use, commercial licence for industrial use) as a future refinement.

21. Final Framework Synthesis (Revised)

21.1 Final Framework Summary

Version 2.0 of the Master Research Framework is an eight-layer, open-source, computational-experimental workflow for rational polymer selection and QbD-driven formulation development of spray-dried amorphous solid dispersions, demonstrated on indomethacin. The framework integrates HSP, Flory-Huggins χ , RDKit descriptors, Gordon-Taylor T_g , and literature evidence into a Composite Compatibility Index with PCA pre-processing; uses multi-expert AHP-TOPSIS for ranking; quantifies uncertainty via joint-distribution Monte Carlo; and defines a logistic-regression-based failure boundary map. The framework is at the end of the framework-development stage and the beginning of the framework-validation stage; generalization to other BCS Class II drugs is future work.

21.2 Advantages (Retained)

Transparent, auditable decision trail aligned with ICH Q8 and Q9; open-source, reproducible, deployable without licensing; approximately 80 to 96 percent reduction in experimental polymer screening burden (revised from 'approximately 96 percent' to 'approximately 80 to 96 percent' to reflect the uncertainty); explicit uncertainty quantification and sensitivity analysis including Morris elementary effects; scientifically grounded in established methods; modular architecture enabling extension; implementable within a 40-week extended M.Pharm timeline. The PCA pre-processing and multi-expert AHP are added advantages in V2.0.

21.3 Limitations (Revised)

Validated on a single drug (indomethacin); generalization is a hypothesis for future work, not a demonstrated result. No bioperformance (dissolution, PBPK) prediction. No scale-up or processability prediction. Lindvig χ estimation has known limitations; uncertainty propagated but bias not fully characterised without MPD calibration. Cross-correlation between HSP and χ is addressed by PCA but not eliminated. AHP weights are subjective; multi-expert elicitation reduces but does not eliminate subjectivity. Binary (drug-polymer) only; ternary not supported. Literature-evidence score is binary (1.0, 0.5, 0.0). Six-month accelerated stability is the ICH Q1A minimum; 12-month long-term is future work. Single-batch pseudoreplication for non-rank-1 polymers (three batches for rank-1 only).

21.4 Novelty Statement (Revised: 3 Claims Retained, 3 Removed)

Three novelty claims are retained: (N2) the Composite Compatibility Index as a transparent integration device with PCA pre-processing (with the explicit caveat that the HSP- χ cross-correlation requires PCA pre-processing and that the CCI may be effectively one-dimensional in some candidate sets); (N3) the operationalisation of failure boundary mapping for indomethacin SD-ASDs via logistic regression (bounded wording, acknowledging prior design-space concepts by Yu 2008 and failure-boundary concepts by Scheirs); (N4) the open-source Python and R release with Docker containerisation and Zenodo DOI archival as the first reproducible computational-experimental framework of this scope for ASD polymer screening.

Three novelty claims from Version 1.0 are removed: (N1) the integrated eight-layer architecture (reclassified as a workflow description, because prior work covers 10 to 12 of the 13 architectural components individually); (N5) the approximately 85 percent material reduction claim (reclassified as an arithmetic projection, because no prospective empirical 25-polymer screen was performed); (N6) the framework transferability claim (reclassified as a hypothesis for future work, because single-drug validation cannot demonstrate generalisability). No claim of a new physical model, equation of state, or thermodynamic theory is made.

21.5 Publication Potential (Revised: Scope-Based Recommendation)

The publication target is revised based on scope rather than impact factor. Version 1.0 named Molecular Pharmaceutics (impact factor approximately 4.6) as the primary target. The committee endorses the reviewer's assessment (criticism 1.4) that this scope is mismatched with a QbD-driven framework paper. The primary submission target is AAPS PharmSciTech (scope fit for QbD and formulation work; acceptance probability 35 to 45%). The secondary target is International Journal of Pharmaceutics (formulation-focused; acceptance probability 18 to 25%). Molecular Pharmaceutics is not recommended unless the computational chemistry depth is substantially increased (e.g., MD validation of χ).

If the experimental programme cannot be completed within the M.Pharm timeline, the framework is reframed as a protocol paper for submission to MethodsX or BMJ Open (acceptance probability 60 to 75%). The graphical abstract specification is updated to a minimum of 800 by 800 pixels (criticism 11.8), the figure count reduced from 10 to 8 for the main text with the remainder moved to supplementary material, and a data availability statement template is added (criticism 11.2).

21.6 Technology Readiness Level (Revised)

TRL 4 (laboratory validation). The framework is validated in a laboratory environment for indomethacin with literature data. TRL 5 (validation in relevant environment) would require prospective wet-lab validation with DSC, FTIR, PXRD, dissolution, and stability data; this is the core of the 40-week roadmap (Section 23). TRL 6 (demonstration in relevant environment) would require pilot-scale spray drying and GMP documentation; out of M.Pharm scope. TRL 7-9 (system prototype, complete system, actual system proven) would require multi-drug validation and industrial deployment; out of scope.

21.7 Industrial Readiness Level (Retained)

IRL 2 to 3 (validated in a laboratory setting but not in an industrial pipeline). IRL 4 to 5 would require the enhancements listed in Section 20.3 (bioperformance prediction, expanded polymer library, ternary prediction, ELN integration, scale-up prediction, Docker-based long-term reproducibility, GMP documentation, CSV, CoG model, FTO analysis, comparison to COSMO-RS). All out of M.Pharm scope.

21.8 Scientific Readiness Level (Revised Downward)

SRL 5 to 6 (revised downward from Version 1.0's SRL 7 to 8). The framework is scientifically rigorous in structure (established methods, transparent assumptions, explicit uncertainty quantification, pre-registered statistical decision rules) but is limited by the single-drug validation and the absence of empirical data. The committee endorses the reviewer's implicit position that SRL 7 to 8 was inflated given the validation gaps. SRL 7 (consensus scientific method) would require multi-drug, multi-laboratory validation and adoption by the broader pharmaceutics community; SRL 9 (consensus scientific method with broad adoption) is out of scope.

21.9 Future Expansion (Retained with Reordered Priorities)

Eight future expansion items are retained, with reordered priorities reflecting V2.0: (1) prospective wet-lab validation on indomethacin (40-week roadmap, highest priority); (2) computational expansion to n equals 20 polymers (literature-only, Phase 2); (3) multi-drug validation on 5 to 10 BCS Class II drugs (felodipine, itraconazole, naproxen, carbamazepine, fenofibrate); (4) bioperformance prediction (Noyes-Whitney plus spring-and-parachute plus PK-Sim PBPK); (5) ternary systems (drug-polymer-surfactant); (6) process prediction (spray-drying process model coupling); (7) COSMO-RS integration (replace Lindvig χ); (8) MD validation (optional layer for selected pairs). The 12-month long-term stability is added as a ninth future item, initiated in parallel with the M.Pharm programme.

22. Risk Register (New in V2.0)

22.0 Risk Classification and Action Thresholds

Risks are classified into five categories: scientific, experimental, statistical, publication, and implementation. Each risk is assigned a likelihood (L, 1 to 5 scale) and impact (I, 1 to 5 scale); the Risk Priority Number (RPN) is L times I. Action thresholds per ICH Q9 (criticism 10.2): RPN less than 15 is acceptable (monitor only); RPN 15 to 30 requires mitigation; RPN greater than 30 requires design change. The risk register is reviewed at each gate (Section 5.2) and a risk dashboard is maintained in the supplementary material.

Table 22.1 — Risk Register (V2.0)

ID	Risk	Category	L	I	RPN	Action	Mitigation	Gate
S1	Cascading hypothesis failure (H1 fails, propagates to H2-H4)	Scientific	3	5	15	Mitigation	Contingency manuscript structure (Section 3.7); independent H testing	G3
S2	FBM statistically indefensible (boundary points insufficient)	Scientific	3	4	12	Acceptable	Minimum 20-30 points; bootstrap CI on boundary	G4
S3	Lindvig chi bias for H-bonded in domethacin	Scientific	4	3	12	Acceptable	MPD calibration for one pair; chi uncertainty plus or minus 25%	G3
S4	Gordon-Taylor or K bias from crystalline density	Scientific	4	3	12	Acceptable	Measure amorphous density; plus or minus 5 K systematic flag	G3
E1	Single-batch pseudoreplication for non-rank-1 polymers	Experimental	4	4	16	Mitigation	3 batches for rank-1; sensitivity analysis for others	G3
E2	PXRD coverage insufficient for top-2	Experimental	2	4	8	Acceptable	PXRD for all top-2 at T=0 (mandatory change 4)	G3
E3	Stability duration insufficient (6-month only)	Experimental	4	4	16	Mitigation	12-month long-term initiated in parallel; future work	G3
E4	Dissolution pH 6.8 only (insufficient biopharmaceutical picture)	Experimental	2	3	6	Acceptable	Multi-pH (1.2, 4.5, 6.8) mandatory in V2.0	G3
E5	Spray-drying yield variability	Experimental	3	3	9	Acceptable	Record yield per batch; report in Decision Report	G3
E6	Tablet compression failure (capping, lamination)	Experimental	3	3	9	Acceptable	Adjust compression force; document in DoE	G3

ID	Risk	Category	L	I	RPN	Action	Mitigation	Gate
ST1	n=6-8 under powered for confirmatory H1	Statistical	5	4	20	Mitigation	Treat as exploratory; expand to n=20 (Phase 2)	G3
ST2	4 pure-error df insufficient for LOF test	Statistical	3	3	9	Acceptable	5 centre points for 6 pure-error df (V2.0)	G3
ST3	Sobol indices unreliable at n=8	Statistical	4	2	8	Acceptable	Removed from primary analysis; Morris instead	G3
ST4	Multiple comparisons family ambiguous	Statistical	3	3	9	Acceptable	Family specified per analysis (Section 15.3)	G3
ST5	Bayesian exclusion unjustified	Statistical	2	3	6	Acceptable	Bayesian sensitivity added (Section 15.7)	G3
P1	Scope mismatch with target journal	Publication	3	4	12	Acceptable	AAPS Pharm SciTech primary; IJP secondary (V2.0)	Submission
P2	Figure/table density too high	Publication	2	3	6	Acceptable	Reduced from 10 to 8 figures (V2.0)	Submission
P3	Graphical abstract below journal minimum	Publication	2	2	4	Acceptable	Updated to 800x800 px minimum (V2.0)	Submission
P4	Closed 70-paper evidence base	Publication	3	3	9	Acceptable	Deposit xlsx on Zenodo with DOI (future)	Submission
P5	No external pre-submission review	Publication	3	3	9	Acceptable	External review mandatory in V2.0 (Section 19.5)	Submission
P6	PRISMA checklist not appended	Publication	2	2	4	Acceptable	Append PRISMA 2020 checklist (future)	Submission
I1	Timeline slippage (40-week roadmap)	Implementation	4	4	16	Mitigation	Phase-by-phase tracking; buffer weeks built in	All

ID	Risk	Category	L	I	RPN	Action	Mitigation	Gate
I2	Multi-expert elicitation infeasible	Implementation	3	3	9	Acceptable	Single-expert fallback with explicit flag	G2
I3	M.Pharm scope boundary (12-month stability infeasible)	Implementation	5	3	15	Mitigation	Initiate in parallel; report in follow-up	G3
I4	Docker container build failure	Implementation	2	3	6	Acceptable	Pin versions; test build in CI	Submission
I5	Software dependency drift over time	Implementation	3	3	9	Acceptable	Docker with pinned versions; Zenodo DOI	Submission

23. Implementation Roadmap (Revised)

23.1 Eight-Phase Plan (40 Weeks, Stopping at Submission)

The implementation roadmap is an eight-phase plan spanning 40 weeks and ending at manuscript submission. No post-publication planning is included (per Task 12 of the user brief). The phases are sequential with some parallelism in Phase 5 (stability initiation).

Table 23.1 — Implementation Roadmap (40 Weeks, V2.0)

Phase	Weeks	Activity	Output	Gate
1	1-4	Computational implementation: 8-layer architecture, PCA, FBM logistic regression	Working Python pipeline	G1 (HSP filter)
2	5-8	Multi-expert AHP elicitation; baseline comparison; LOO CV; expand to n=20 (literature)	AHP weights; ranking; CV results	G2 (CR<0.08)
3	9-16	Spray drying DoE (3 batches for rank-1); PXRD for top-2; GC-HS; HPLC-UV Q2(R2)	Spray-dried ASDs; analytical data	G3 (validation)
4	17-20	Tablet compression DoE (5 centre points); multi-pH dissolution (1.2, 4.5, 6.8)	Tablets; dissolution profiles	G3
5	21-24	6-month accelerated stability initiation (40C/75%RH); 12-month long-term initiation (parallel)	Stability study underway	G4 (FBM pre-reg)
6	25-32	Statistical analysis; Bayesian sensitivity; equivalence testing (TOST); FBM estimation	Statistical report; FBM contour	G3, G4
7	33-36	Manuscript drafting; figures; supplementary; external expert review	Draft manuscript; external review	Submission prep
8	37-40	Revision; submission to AAPS PharmSciTech (primary) or IJP (secondary)	Submitted manuscript	Submission

23.2 Phase Dependencies and Critical Path

The critical path runs through Phases 1, 3, 5, 6, and 8. Phase 2 (multi-expert AHP) can be partially parallelised with Phase 1 (computational implementation). Phase 4 (tablet compression) depends on Phase 3 (spray-dried ASDs). Phase 5 (stability) is initiated at the end of Phase 3 and runs in parallel with Phases 4 and 6; the 6-month stability data are available at week 32 (week 24 plus 6 months equals approximately 50 weeks, but the framework reports 3-month interim data at week 32 and 6-month data at week 50 if available). Phase 7 (manuscript drafting) depends on Phase 6 (statistical analysis). Phase 8 (submission) depends on Phase 7 and the external review.

If the 6-month stability data are not available by week 36 (manuscript drafting), the manuscript reports 3-month interim data with explicit acknowledgment that 6-month data will be reported in a follow-up. This is a known risk (I3 in the risk register) with mitigation.

23.3 Resource Requirements

Computational resources: standard laptop (8 GB RAM, Python 3.11+, Docker). No GPU or cluster required. Software: open-source only (RDKit, numpy, pandas, scipy, matplotlib, scikit-learn, SALib, PyMC or Stan for Bayesian). Wet-lab resources: Buchi B-290 spray dryer, DSC (TA Instruments or equivalent), FTIR with ATR accessory, PXRD (Bruker D8 or equivalent), HPLC-UV, GC-HS, SEM (optional), laser diffraction (optional), dissolution apparatus (USP II), Karl Fischer titrator, stability chambers (25C/60%RH and 40C/75%RH). Materials: indomethacin (gamma-form, 25 g), six polymers (5 g each), solvents (acetone, 5 L), tablet excipients (MCC, croscarmellose, Mg stearate). Personnel: M.Pharm student (full-time), supervisor (part-time), three expert panellists (1 hour each for AHP elicitation), one external reviewer (4 hours for pre-submission review).

24. References (Updated)

24.1 References (Retained from V1.0 plus 10 New)

The 67 references from Version 1.0 are retained. Ten new references are added for V2.0: (68) Hosseini et al. (2022) on PCA in pharmaceutical formulation; (69) Acuna et al.. (2020) on multi-expert AHP elicitation protocols; (70) Morris (1991) on elementary effects screening; (71) Saltelli et al. (2010) on variance-based sensitivity analysis; (72) FDA (2010) on Bayesian stability analysis acceptance; (73) ICH Q2(R2) (2023) on analytical method validation; (74) ICH Q3C (2021) on residual solvents; (75) Bland and Altman (1986) on equivalence testing (TOST); (76) AAPS PharmSciTech (2024) on journal scope; (77) Yu et al. (2014) on logistic regression for design-space boundary estimation. The full numbered reference list (Vancouver style) is provided in the supplementary material.

Appendix A — Review Matrix

This appendix maps every distinct criticism from the Critical Scientific Review (CSR-2026-IND-ASD-001) to the corresponding section of the Master Framework, the committee's classification (Mandatory, Recommended, Optional, Future Work, Unsupported Reviewer Opinion), and the disposition (Implemented, Modified, Removed, Deferred). The matrix is the master traceability artifact.

Table A.1 — Review Matrix (Criticism to Disposition Mapping)

Criticism ID	Framework Section	Criticism Summary	Committee Classification	Disposition	V2.0 Section
1.1(1)	Whole framework	Framework is architecture, not validation; desk-reject risk	Mandatory	Implemented	Exec Summary, Sec 11.0
1.1(2)	CompFW validation set	n=8 statistically underpowered	Mandatory	Implemented	Sec 3.3, 15.2
1.1(3)	Hypothesis architecture	Cascading-failure fragility	Mandatory	Implemented	Sec 3.2-3.7
1.1(4)	Publication target	Scope mismatch with Mol Pharm	Mandatory	Implemented	Sec 21.5
1.1(5)	Novelty claims	N1, N5, N6 inflated	Mandatory	Implemented	Sec 1.4, 21.4
N1	Novelty claim N1	8-layer architecture not novel (prior art covers 10-12 components)	Mandatory	Removed	Sec 1.4
N2	Novelty claim N2	CCI novel but double-counts cohesive energy	Recommended	Modified	Sec 4.3, 7.1-7.7
N3	Novelty claim N3	FBM operationalisation only; concept pre-exists	Recommended	Modified	Sec 18
N4	Novelty claim N4	Open-source is genuine novelty	Optional	Retained	Sec 21.4
N5	Novelty claim N5	85% reduction is arithmetic, not a finding	Mandatory	Removed	Sec 1.4
N6	Novelty claim N6	Transferability is hypothesis, not novelty	Mandatory	Removed	Sec 1.4, 3.6
H1	Hypothesis H1	Weakly answerable at n=8	Mandatory	Modified	Sec 3.3
H2	Hypothesis H2	Answerable (RMSE <= 10 K)	Optional	Retained	Sec 3.4
H3	Hypothesis H3	Weakly answerable (n=12 vessels marginal)	Recommended	Modified	Sec 3.5
H4	Hypothesis H4	Weakly answerable; no fallback boundary	Mandatory	Implemented	Sec 3.5, 18
H5	Hypothesis H5	Non-falsifiable (literature-anchored)	Mandatory	Removed	Sec 3.6
3.1a	H1 operationalisation	Conflates ranking with dissolution	Mandatory	Implemented	Sec 3.1
3.1b	H4 dependency	No fallback empirical boundary	Mandatory	Implemented	Sec 3.5, 18
3.1c	H5 falsifiability	Non-falsifiable	Mandatory	Implemented	Sec 3.6
3.2	Statistical decision rules	Missing pre-registration elements	Mandatory	Implemented	Sec 15.0
4.1	HSP-chi cross-correlation	Acknowledged not resolved	Mandatory	Implemented	Sec 4.3, 6.2
4.2	AHP single-expert	Incompatible with ICH Q9	Mandatory	Implemented	Sec 4.4, 8.4
4.3	Lindvig chi additivity	Invalid for H-bonded systems	Recommended	Modified	Sec 2.2, 16.6
4.4	Simha-Boyer density	Uses crystalline density	Recommended	Modified	Sec 2.3, 17.1
4.5	CCI correlated inputs	Effective dimensionality unknown	Mandatory	Implemented	Sec 4.3, 7.7
4.6	Monte Carlo UQ	Sampling, not UQ	Recommended	Modified	Sec 10.2
4.7	No held-out test set	In-sample validation	Mandatory	Implemented	Sec 11.2, 17.7-17.8
5.1	DSC+FTIR insufficiency	AAPS undetectable at <20 nm	Mandatory	Implemented	Sec 16.1
5.2	PXRD coverage	Top-2 not characterised	Mandatory	Implemented	Sec 16.1
5.3	Single-batch pseudoreplication	Mixed-effects degenerate	Mandatory	Implemented	Sec 16.2, 15.9

Criticism ID	Framework Section	Criticism Summary	Committee Classification	Disposition	V2.0 Section
5.4	Dissolution pH 6.8 only	Ignores pH-dependent solubility	Recommended	Implemented	Sec 16.4
5.5	Stability duration	6-month is ICH Q1A minimum, not Q1 standard	Mandatory	Implemented	Sec 16.5
5.6	Missing analytical methods	GC-HS, SEM/PSA, HPLC-UV Q2(R2)	Mandatory	Implemented	Sec 16.6
6.1	Spearman rho at n=8	Cannot confirm or refute H1	Mandatory	Implemented	Sec 3.3, 15.2
6.2	ANOVA 4 pure-error df	Lacks power for quadratic terms	Recommended	Modified	Sec 16.2-16.3
6.3	Bootstrap CI for f2	Suboptimal; permutation test preferred	Optional	Implemented	Sec 15.5
6.4	Multiplicity family unclear	Family not specified	Recommended	Implemented	Sec 15.3
6.5	R2_pred threshold lenient	0.50 too low; should be 0.70	Recommended	Implemented	Sec 15.6
6.6	No equivalence testing for FBM	Binary classification, not equivalence	Mandatory	Implemented	Sec 11.5, 15.8
6.7	Bayesian exclusion unjustified	ICH Q1E does not preclude Bayesian	Recommended	Implemented	Sec 15.7
6.8	Sobol indices unreliable at n=8	88 evaluations insufficient	Recommended	Removed	Sec 9.4
6.9	Pseudoreplication mitigation is admission	Not solution	Mandatory	Implemented	Sec 15.9
7.1	Circular validation	Predictions validated against experiments designed using predictions	Mandatory	Implemented	Sec 11.1
7.2	No independent test set	In-sample validation	Mandatory	Implemented	Sec 11.2
7.3	Success/Partial/Fail arbitrary	50% threshold unjustified	Recommended	Modified	Sec 11.3
7.4	H5 non-falsifiable	Cannot be falsified in scope	Mandatory	Implemented	Sec 3.6
7.5	FBM lacks formal definition	Operational threshold, not statistical	Mandatory	Implemented	Sec 18
7.6	5-step closed loop procedural	Not statistical	Recommended	Modified	Sec 11.4
7.7	No baseline comparison	Cannot demonstrate marginal value	Recommended	Implemented	Sec 11.4, 17.6
8.1	FBM taxonomy no derivation	Warning region undefined	Mandatory	Implemented	Sec 18.2
8.2	2D contour ignores 2 factors	Slice through 4D space	Recommended	Implemented	Sec 18.3
8.3	~5 boundary points indefensible	Minimum 20-30 required	Mandatory	Implemented	Sec 18.4
8.4	No comparison to design-space edge	FBM vs ICH Q8 design space	Recommended	Implemented	Sec 18.8
8.5	Failure terminology misleading	Operational, not safety	Recommended	Implemented	Sec 18.7
8.6	Cherry-picking risk	Post-hoc boundary definition	Mandatory	Implemented	Sec 18.5
8.7	No UQ on boundary location	Deterministic contour	Recommended	Implemented	Sec 18.6

Criticism ID	Framework Section	Criticism Summary	Committee Classification	Disposition	V2.0 Section
9.1-9.9	Industrial adoption	TRL 4, no GMP, no scale-up, no CoG, no FTO, no COSMO-RS comparison, IP risk, no regulatory pathway	Future Work	Deferred	Sec 20, 21.9
10.1	Q8 design space as research output	Not regulatory deliverable	Future Work	Deferred	Sec 20
10.2	Q9 FMEA without RPN thresholds	Action thresholds not defined	Mandatory	Implemented	Sec 22.0
10.3	Q10 PQS elements missing	Aspirational not actual	Future Work	Deferred	Sec 2.6, 20
10.4	ICH Q2(R2) no method validation	Analytical methods not validated	Mandatory	Implemented	Sec 16.6
10.5	ICH Q3C no GC-HS method	Residual solvent unverifiable	Mandatory	Implemented	Sec 16.6
10.6	ICH Q1A 6-month only	Not 12-month long-term	Mandatory	Implemented	Sec 16.5
10.7	ICH Q1E no statistical procedure	Trend evaluation unspecified	Recommended	Implemented	Sec 16.5
10.8	Zenodo not regulatory pre-submission	Misrepresentation	Recommended	Modified	Sec 18.5
10.9	ICH M9 biowaiver implications	Not discussed	Optional	Deferred	Future work
11.1	Figure/table density too high	15+15 in 8-10k words	Recommended	Implemented	Sec 13.1
11.2	No data availability template	Required by journals	Recommended	Implemented	Sec 21.5
11.3	Supplementary modules unspecified	Size/scope unclear	Optional	Deferred	Supplementary
11.4	Docker no DOI	Not deposited	Recommended	Implemented	Sec 14.4
11.5	70-paper evidence base closed	xlsx not public	Recommended	Deferred	Future work
11.6	No OSF pre-registration of code	Post-hoc analysis risk	Optional	Deferred	Future work
11.7	PRISMA checklist missing	Compliance unverifiable	Optional	Deferred	Future work
11.8	Graphical abstract too small	530x530 below 800x800 minimum	Mandatory	Implemented	Sec 13.2
11.9	Reviewer simulation internal	No external review	Mandatory	Implemented	Sec 19.5
11.10	Citation counts approximate	CrossRef not verified	Optional	Deferred	Future work
Q-MAJ-001 to 040	Various	40 major reviewer questions	Mandatory/Recommended	Implemented/Modified	Throughout
Q-MOD-001 to 035	Various	35 moderate reviewer questions	Recommended/Optional	Implemented/Deferred	Throughout
Q-MIN-001 to 025	Various	25 minor reviewer questions	Optional	Implemented/Deferred	Throughout
Flaws 1-15	Various	15 ranked fatal flaws	Mandatory/Recommended/Optional	Implemented/Modified/Removed	Throughout

Criticism ID	Framework Section	Criticism Summary	Committee Classification	Disposition	V2.0 Section
14.1 Priority 1	Experimental programme	Execute Phases 3-6	Mandatory	Implemented	Sec 23
14.1 Priority 2	Computational validation	Expand to $n \geq 20$	Mandatory	Implemented	Sec 23, Phase 2
14.1 Priority 3	FBM logistic regression	Define FBM statistically	Mandatory	Implemented	Sec 18
14.1 Priority 4	Negative controls	Include rejected polymers	Mandatory	Implemented	Sec 11.1
14.1 Priority 5	Multi-expert AHP	Geometric-mean primary	Mandatory	Implemented	Sec 4.4, 8.4

Appendix B — Change Log

This appendix provides per-revision traceability from Version 1.0 to Version 2.0. Each entry records the original text or claim, the revised text or claim, the reason for the change, the reviewer criticism addressed, the scientific justification, and the classification tag (Retained, Modified, Removed, Deferred).

Table B.1 — Change Log (Per-Revision Entries)

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-001	Abstract / Exec Summary	'reduce experimental polymer screening by approximately 80-96%'	'approximately 80-96% reduction in experimental polymer screening burden (revised range to reflect uncertainty)'	Soften specific quantitative claim to range	1.1(5), N5	85% was arithmetic projection, not validated finding	Modified
R-002	Abstract / Exec Summary	'predict Tg within approximately 10 K'	'predict Tg within approximately 10 K (RMSE acceptance criterion; exploratory at n=6-8, confirmatory at n>=20)'	Add exploratory /confirmatory classification	6.1	n=6-8 CI too wide for confirmatory claim	Modified
R-003	Sec 1.3 Objective 8	'Be of publication quality suitable for submission to a Q1 pharmaceuticals journal'	'Be of publication quality suitable for submission to a Q1 pharmaceuticals journal AFTER empirical validation is completed'	Add 'after empirical validation' qualifier	1.1(1)	Framework is plan, not completed study	Modified
R-004	Sec 1.4 Novelty	6 novelty claims (N1-N6)	3 novelty claims retained (N2, N3, N4); 3 removed (N1, N5, N6)	Remove inflated claims	1.1(5), 2.1, N1, N5, N6	N1 is workflow not contribution; N5 is arithmetic; N6 is hypothesis	Removed
R-005	Sec 1.4 Novelty Declaration	'No new equation of state, no new thermodynamic model, no machine learning'	'No new equation of state, no new thermodynamic model, no machine learning. Three retained novelty claims: CCI with PCA caveat, FBM operationalisation (bounded), open-source release.'	Explicitly list retained claims	2.1, 2.2	Bounded novelty is defensible	Modified
R-006	Sec 3 (NEW)	Implicit hypotheses only	Explicit H1-H4 with primary/secondary classification; H5 dropped	Make hypotheses testable	3.1a-c, 3.2	H5 non-falsifiable; H1-H4 need pre-registered decision rules	Modified
R-007	Sec 3.3 H1	'rho >= 0.70 acceptance'	'rho >= 0.70; exploratory at n=6-8, confirmatory at n>=20'	Add sample-size caveat	6.1	CI at n=6-8 too wide	Modified

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-008	Sec 3.6 H5	H5 as formal hypothesis	H5 dropped; reframed as Discussion topic	Drop non-falsifiable hypothesis	3.1c, 7.4	Literature-anchored validation cannot falsify	Removed
R-009	Sec 3.7 (NEW)	No contingency plan	Contingency manuscript structure if H1 fails	Plan for H1 failure	1.1(3), Q-MAJ-040	Cascading failure needs contingency	Implemented
R-010	Sec 4.3 Layer 5	Direct weighted linear combination CCI	PCA pre-processing mandatory before CCI	Address HSP-chi double-counting	4.1, 4.5	$\rho(\text{HSP}, \text{chi}) = 0.94$; PCA extracts orthogonal info	Modified
R-011	Sec 4.4 Layer 6	Single-expert AHP primary	Multi-expert AHP primary (3-5 experts, geometric-mean)	Address ICH Q9 governance mismatch	4.2, 9.4	Single-expert incompatible with cross-functional input	Modified
R-012	Sec 4.5 Layer 7	FBM as operational 80% threshold	FBM as logistic regression $P(\text{failure}) = 0.5$ contour	Statistical definition of boundary	7.5, 8.1	Operational threshold is binary, not statistical	Modified
R-013	Sec 5.2 Gates	3 gates (HSP, AHP $\text{CR} < 0.10$, validation)	4 gates (added Gate 4 for FBM pre-registration); CR threshold tightened to 0.08	Tighten CR; add FBM gate	Q-MOD-001, 8.6	0.08 stricter; FBM must be pre-registered	Modified
R-014	Sec 6.2 Eq 10	$\text{CCI} = \sum(w_i * s_i)$ on raw scores	$\text{CCI} = \sum(w_j * \text{PC}_j)$ on retained principal components	PCA pre-processing	4.1, 4.5	Removes double-counting	Modified
R-015	Sec 6.3 Eq 11 (NEW)	No FBM equation	$\text{logit}(P(\text{failure})) = \beta_0 + \beta_1 * \text{rank} + \beta_2 * \text{inlet}_T + \beta_3 * \text{loading} + \beta_4 * \text{feed}_{\text{conc}}$	Formal FBM definition	7.5, 8.1	Logistic regression provides probabilistic boundary	Implemented
R-016	Sec 6.4 Eq 12 (NEW)	No boundary uncertainty equation	Bootstrap CI on contour ($n_{\text{boot}} = 10,000$)	UQ on boundary location	8.7	Deterministic contour has no uncertainty estimate	Implemented
R-017	Sec 7.5 Limitations	5 limitations	6 limitations (added PCA interpretation caveat)	Acknowledge PCA limitation	4.5	PCs may lack direct physical interpretation	Modified
R-018	Sec 7.6 Validation	Spearman $\rho \geq 0.7$ against literature	Spearman $\rho \geq 0.7$; exploratory at $n=6-8$; LOO CV added; held-out test set added	Add CV and held-out protocols	4.7, 7.2	In-sample validation insufficient	Modified
R-019	Sec 7.7 (NEW)	No effective dimensionality reporting	Mandatory PCA effective dimensionality reporting	Address unknown informational content	4.5	Effective dimensionality must be transparent	Implemented

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-020	Sec 8.4 (NEW)	No multi-expert protocol	7-step multi-expert elicitation protocol with Kendall W	Specify multi-expert method	4.2, 9.4	Geometric-mean aggregation; inter-expert agreement	Implemented
R-021	Sec 9.1 Sensitivity	OAT + Monte Carlo	OAT + Monte Carlo + Morris elementary effects	Detect interaction effects	Q-MOD-002	OAT cannot detect interactions	Modified
R-022	Sec 9.4 (NEW)	Sobol indices in primary analysis	Sobol removed from primary; relegated to future work	Address underpowered Sobol	6.8	N=8 insufficient for Sobol (need 100-1000)	Removed
R-023	Sec 10.2 UQ	Monte Carlo N=10,000 with +20% uniform	Joint distribution propagation of all 7 uncertainty sources	True UQ not sampling	4.6	Joint distribution required for UQ claim	Modified
R-024	Sec 10.5 (NEW)	No convergence criterion	R-hat < 1.01 across 5 chains OR CoV < 5% between N=10k and N=50k	Specify convergence	Q-MOD-021	Convergence must be verified	Implemented
R-025	Sec 11.0 (NEW)	No development/validation/generalization distinction	Explicit three-stage distinction	Clarify framework stage	7.1-7.7	Conflation of stages	Implemented
R-026	Sec 11.1 (NEW)	No negative controls	1-2 rejected polymers included in wet-lab	Break circularity	7.1	Test screen's specificity	Implemented
R-027	Sec 11.2 (NEW)	No held-out test set	LOO CV + held-out 2 polymers	External validation	4.7, 7.2	In-sample validation insufficient	Implemented
R-028	Sec 11.4 (NEW)	No baseline comparison	HSP-only and equal-weight baselines; delta rho >= 0.10	Demonstrate marginal value	7.7	Complexity must be justified	Implemented
R-029	Sec 11.5 (NEW)	No equivalence testing	TOST with +-5% margin for FBM boundary	Statistical boundary confirmation	6.6	Binary threshold not equivalence	Implemented
R-030	Sec 13.1 Figures	10 figures	8 main-text figures (2 moved to supplementary); 2 new (PCA scree, FBM contour)	Reduce density; add new visualizations	11.1	Density too high; new methods need figures	Modified
R-031	Sec 13.2 Graphical abstract	530x530 px	800x800 px minimum (1320x1320 preferred)	Meet journal minimum	11.8	Below journal specification	Modified
R-032	Sec 14.4 Docker	Referenced but no DOI	Deposited on Docker Hub/GHCR with Zenodo DOI	Reproducibility	11.4	Container must be archived	Modified
R-033	Sec 15.0 (NEW)	No pre-registered decision rules	7-element pre-registration for each H1-H4	Pre-register analysis	3.2	Avoid post-hoc analysis	Implemented

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-034	Sec 15.2 Sample size	n=6-8 acceptable	n=6-8 exploratory; n>=20 confirmatory	Acknowledge underpowering	6.1	CI too wide at n=6-8	Modified
R-035	Sec 15.3 Multiple comparisons	Tukey/Dunnett/ BH-FDR specified	Family explicitly specified per analysis	Specify multiplicity family	6.4	Family was ambiguous	Modified
R-036	Sec 15.6 R2_pred	>= 0.50	>= 0.70	Tighten threshold	6.5	0.50 too lenient for Q1	Modified
R-037	Sec 15.7 (NEW)	Bayesian excluded	Bayesian sensitivity analysis added	Address unjustified exclusion	6.7	ICH Q1E does not preclude Bayesian	Implemented
R-038	Sec 15.8 (NEW)	No equivalence testing framework	TOST framework for FBM	Statistical boundary	6.6	Binary threshold insufficient	Implemented
R-039	Sec 15.9 (NEW)	Mixed-effects model as mitigation	3 batches for rank-1; sensitivity analysis fallback	Address degenerate random effect	5.3, 6.9	Single-batch mixed-effects is degenerate	Modified
R-040	Sec 16 (NEW)	No dedicated experimental section	New Section 16 consolidating experimental methods	Consolidate experimental revisions	5.1-5.6, E1-E9	Experimental design needs dedicated treatment	Implemented
R-041	Sec 16.1 Characterization	DSC + FTIR + PXRD (final only)	DSC (MDSC option) + FTIR (ATR) + PXRD for ALL top-2 at T=0	Address PXRD coverage gap	5.1, 5.2, E1, E2	AAPS undetectable; top-2 must be characterised	Modified
R-042	Sec 16.2 Spray drying	Single-batch Box-Behnken	3 batches for rank-1; 5 centre points (was 3)	Address pseudo replication + LOF power	5.3, 6.2, E3	Single-batch degenerate; 4 df insufficient	Modified
R-043	Sec 16.3 Tablet	2^3 factorial, 3 centre points	CCD with 5 centre points + 4 axial	Better curvature detection	Q-MOD-004, E5	3 centre points underpowered	Modified
R-044	Sec 16.4 Dissolution	pH 6.8 only	Multi-pH (1.2, 4.5, 6.8) mandatory	Biopharmaceutical completeness	5.4, E4	Indomethacin pKa 4.5; pH-dependent solubility	Modified
R-045	Sec 16.5 Stability	6-month accelerated only	6-month accelerated + 12-month long-term (future work, initiated in parallel)	Q1 publication standard	5.5, E6, 10.6	6-month is ICH Q1A minimum, not Q1 standard	Modified
R-046	Sec 16.6 Analytical (NEW)	No GC-MS, SEM/PSA, HPLC-UV Q2(R2)	GC-MS (ICH Q3C) + HPLC-UV Q2(R2) + SEM/PSA (recommended)	Regulatory compliance	5.6, 10.4, 10.5, E7, E8	Methods missing; Q3C/Q2(R2) compliance	Implemented
R-047	Sec 17.5 CCI	Direct weighted combination	PCA + CCI on retained PCs	Apply PCA pre-processing	4.1, 4.5	Double-counting addressed	Modified
R-048	Sec 17.6 (NEW)	No baseline comparison in worked example	HSP-only and equal-weight baselines reported	Demonstrate marginal value	7.7	Complexity must be justified	Implemented
R-049	Sec 17.7-17.8 (NEW)	No LOO CV or held-out test set	LOO CV + held-out 2 polymers	External validation	4.7, 7.2	In-sample validation insufficient	Implemented

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-050	Sec 17.9 Validation	'excellent agreement (rho=0.83)'	'exploratory (rho=0.83; 95% CI 0.20-0.98; requires n>=20 for confirmation)'	Honest CI reporting	6.1	CI too wide for confirmatory claim	Modified
R-051	Sec 18 (NEW)	FBM scattered across sections	Dedicated Section 18 on FBM	Consolidate FBM treatment	7.5, 8.1-8.7	FBM is primary novelty; needs dedicated section	Implemented
R-052	Sec 18.1 FBM definition	Operational 80% threshold	Logistic regression P(failure)=0.5 contour	Statistical definition	7.5, 8.1	Operational threshold is not statistical	Modified
R-053	Sec 18.2 Taxonomy	Safe/Warning/Failure; Warning undefined	Safe (<0.30), Warning (0.30-0.70), Failure (>0.70)	Quantify Warning region	8.1	Warning was undefined	Modified
R-054	Sec 18.3 Multidimensional	2D contour ignores 2 factors	4D with 2D slices at multiple combinations	Address dimensionality reduction	8.2	2D slice misses 4D structure	Modified
R-055	Sec 18.4 Estimation	~5 boundary points	Minimum 20-30 points	Statistical defensibility	8.3	5 points indefensible for contour	Modified
R-056	Sec 18.5 Pre-registration	No boundary pre-registration	Pre-registered on Zenodo before wet-lab	Prevent cherry-picking	8.6	Post-hoc boundary is cherry-picking	Implemented
R-057	Sec 18.6 UQ on boundary	Deterministic contour	Bootstrap CI (n _{boot} =10,000)	Quantify boundary uncertainty	8.7	Deterministic contour has no UQ	Implemented
R-058	Sec 18.7 Terminology	'failure'	'operational failure' or 'CQA failure'	Avoid overstating consequences	8.5	Q30=79% vs 81% not meaningful difference	Modified
R-059	Sec 18.8 Comparison	No comparison to ICH Q8 design space	Explicit comparison	Acknowledge prior concept	8.4	FBM is operationalisation, not new concept	Implemented
R-060	Sec 19.1-19.3 Reviewer simulation	IJP, Mol Pharm, EJPB	AAPS PharmSciTech (primary), IJP (secondary), Mol Pharm (tertiary, scope-mismatch)	Scope-based target	1.4	Mol Pharm scope mismatch	Modified
R-061	Sec 19.5 (NEW)	Internal review only	External pre-submission review mandatory	Genuine external check	11.9	Internal review insufficient	Implemented
R-062	Sec 20 Industrial	IRL 2-3; Tier-1 pre-IND	IRL 2-3 (retained); academic contribution vs industrial applicability vs future validation distinguished	Honest industrial positioning	9.1-9.9	Framework not industrial-ready	Modified
R-063	Sec 21.4 Novelty	6 claims	3 claims retained (N2, N3, N4)	Remove inflated claims	1.1(5), 2.1	N1, N5, N6 not defensible	Modified

Rev ID	Section (V1.0 -> V2.0)	Original (V1.0)	Revised (V2.0)	Reason	Reviewer Criticism	Scientific Justification	Classification
R-064	Sec 21.5 Publication	Mol Pharm primary (IF 4.6)	AAPS PharmSciTech primary (scope fit); IJP secondary	Scope over IF	1.4	Mol Pharm scope mismatch	Modified
R-065	Sec 21.8 SRL	SRL 7-8	SRL 5-6	Honest readiness assessment	16 (dimension 10)	SRL 7-8 inflated given validation gaps	Modified
R-066	Sec 22 (NEW)	No formal risk register	25 risks across 5 categories with RPN and action thresholds	Formal risk management	10.2	RPN thresholds required by ICH Q9	Implemented
R-067	Sec 23 (NEW)	36-week roadmap	40-week 8-phase roadmap stopping at submission	Revised execution plan	14.1	Roadmap must end at submission	Modified
R-068	Sec 2.7 (NEW)	No comparison to in-silico tools	Section 2.7 comparing to COSMO-RS, GastroPlus, Simcyp	Industrial positioning	9.7	Industry uses these tools; framework must position	Implemented
R-069	Sec 2.6 QbD	Q10 alignment claimed without qualification	Q10 alignment aspirational; PQS elements out of scope	Honest regulatory alignment	10.3	PQS is industrial-scale	Modified
R-070	Appendix A (NEW)	No review matrix	Complete review matrix mapping all criticisms	Traceability	Task 1	Mandatory traceability artifact	Implemented
R-071	Appendix B (NEW)	No change log	Per-revision change log	Traceability	Task 13	Mandatory audit trail	Implemented
R-072	Appendix C (NEW)	No rescored audit	10-dimension rescored audit with V1.0 vs V2.0 comparison	Quality assessment	Task 15	Mandatory quality audit	Implemented
R-073	Appendix D (NEW)	No freeze statement	Formal freeze statement	Framework freezing	Task 16	Mandatory freeze before implementation	Implemented
R-074	Throughout	'first operationalised failure boundary map'	'first operationalised failure boundary map for indomethacin SD-ASDs (bounded)'	Bounded novelty wording	1.1(5), N3	Concept pre-exists (Yu, Scheirs)	Modified
R-075	Throughout	'approximately 96% reduction'	'approximately 80-96% reduction (range to reflect uncertainty)'	Soften specific claim	N5	96% was specific; range is honest	Modified
R-076	Throughout	'framework is scientifically generalizable'	'framework is designed to be scientifically transferable; generalization is future work'	Reframe generalization as future work	N6, H5	Single-drug validation cannot demonstrate	Modified

Appendix C — Rescored Quality Audit

This appendix presents the 10-dimension quality audit for Version 2.0, compared against the Version 1.0 self-score and the Critical Scientific Review panel score. For each dimension, the V1.0 score, V2.0 score, and delta are reported, with a narrative explanation of what changed and why.

Table C.1 — Rescored Quality Audit (V1.0 vs V2.0)

#	Dimension	V1.0 Self	Panel Score	V2.0 Committee	Delta (V2.0-V1.0)	Narrative
1	Scientific Quality	82	55	68	-14	V2.0 implements PCA, multi-expert AHP, pre-registered decision rules, and contingency planning. Score decreases from V1.0 self-assessment because V1.0 was inflated; increases from panel score because revisions address the cascading-failure and validation-circularity concerns. Honest positioning: rigorous in structure, limited by single-drug scope.
2	Novelty	76	62	58	-18	V2.0 removes 3 of 6 novelty claims (N1 architecture, N5 85% reduction, N6 transferability). Score decreases from both V1.0 and panel because the committee agrees the removed claims were inflated. The 3 retained claims (CCI with PCA, FBM operationalisation bounded, open-source) are defensible but modest. Honest positioning: bounded novelty, not breakthrough.

#	Dimension	V1.0 Self	Panel Score	V2.0 Committee	Delta (V2.0-V1.0)	Narrative
3	Computational Design	85	68	76	-9	V2.0 adds PCA pre-processing, Morris sensitivity, joint-distribution UQ, convergence criteria, and baseline comparison. Score decreases from V1.0 self because the HSP-chi double-counting is now acknowledged; increases from panel because the PCA and UQ revisions address the double-counting and sampling-not-UQ criticisms. Honest positioning: sound architecture with explicit limitations.
4	Experimental Design	78	48	62	-16	V2.0 adds PXRD for top-2, 3 batches for rank-1, multi-pH dissolution, 5 centre points, GC-HS, HPLC-UV Q2(R2), SEM/PSA. Score decreases from V1.0 self because V1.0 was inflated; increases from panel because the experimental additions address the PXRD, pseudoreplication, pH, and analytical-methods criticisms. Honest positioning: improved but still limited by 6-month stability and single-punch tablet press.

#	Dimension	V1.0 Self	Panel Score	V2.0 Committee	Delta (V2.0-V1.0)	Narrative
5	Statistics	80	50	65	-15	V2.0 adds pre-registered decision rules, Bayesian sensitivity, permutation test, equivalence testing, R2_pred tightened to 0.70, Sobol removed. Score decreases from V1.0 self because the n=6-8 underpowering is now acknowledged; increases from panel because the statistical framework is now more complete. Honest positioning: rigorous in structure, limited by small-n.
6	Industrial Relevance	71	45	48	-23	V2.0 honestly distinguishes academic contribution from industrial applicability and future industrial validation. Score decreases from V1.0 self because the framework is acknowledged as not industrial-ready (TRL 4, no GMP, no scale-up, no CoG, no FTO, no COSMO-RS comparison). Honest positioning: academically defensible, industrially limited.

#	Dimension	V1.0 Self	Panel Score	V2.0 Committee	Delta (V2.0-V1.0)	Narrative
7	Publication Readiness	84	60	64	-20	V2.0 reduces figure count, updates graphical abstract size, adds data availability template, specifies AAPS PharmSciTech as primary target, adds external review. Score decreases from V1.0 self because the scope mismatch with Mol Pharm is acknowledged; increases from panel because the publication revisions address the density, graphical abstract, and external review criticisms. Honest positioning: ready for AAPS PharmSciTech after empirical validation.
8	Q1 Readiness	N/A	35	42	N/A	V2.0 is a research plan, not a completed study; Q1 readiness requires empirical validation. Score increases from panel because the 40-week roadmap specifies the path to Q1 submission. Honest positioning: Q1 readiness achievable in 12-18 months.
9	Technology Readiness (TRL)	42	38	40	-2	V2.0 retains TRL 4 (laboratory validation). Score essentially unchanged; the 40-week roadmap specifies the path to TRL 5 (validation in relevant environment). Honest positioning: TRL 4, path to TRL 5 specified.

#	Dimension	V1.0 Self	Panel Score	V2.0 Committee	Delta (V2.0-V1.0)	Narrative
10	Research Readiness	N/A	40	58	N/A	V2.0 is frozen; no further methodological changes before implementation. Score increases from panel because the freeze statement and the implementation roadmap provide a clear path forward. Honest positioning: research-ready, awaiting execution.

Composite score (committee-weighted): V1.0 self 78/100; panel 51/100; V2.0 committee 69/100. The V2.0 score is higher than the panel score because the revisions address the mandatory and recommended criticisms. The V2.0 score is lower than the V1.0 self-score because the committee removes inflated claims and acknowledges limitations honestly. The committee considers the V2.0 score of 69/100 to be a defensible and honest assessment of the framework's quality as a research plan awaiting empirical validation.

Appendix D — Freeze Statement

This appendix contains the formal freeze statement for Master Research Framework Version 2.0.

Master Research Framework Version 2.0 is hereby FROZEN as of the date of this document. No further methodological changes shall occur before computational implementation begins. Only corrections arising from experimental evidence may modify the framework thereafter.

This freeze is issued by the Scientific Revision Committee, comprising: Editor-in-Chief; Associate Editor; Pharmaceutical Scientist; Computational Pharmaceutics Scientist; Formulation Scientist; QbD Expert; Industrial R&D; Scientist; Regulatory Scientist; Senior Biostatistician.

The freeze statement signifies that the framework's architecture (eight layers), mathematical framework (Equations 1-12), hypothesis structure (H1-H4 with H5 dropped), validation strategy (negative controls, held-out test set, baseline comparison, equivalence testing), statistical framework (pre-registered decision rules, Bayesian sensitivity, permutation test, TOST), experimental framework (PXRD for top-2, 3 batches for rank-1, multi-pH dissolution, 5 centre points, GC-HS, HPLC-UV Q2(R2), SEM/PSA), risk register (25 risks with RPN thresholds), and implementation roadmap (40-week 8-phase plan) are fixed and will not be revised before computational implementation begins.

Corrections arising from experimental evidence may modify the framework thereafter. Such corrections include: (a) revision of AHP weights based on LOO cross-validation results; (b) revision of the Lindvig chi conversion based on MPD calibration data; (c) revision of the Gordon-Taylor K based on measured amorphous density; (d) revision of the FBM logistic regression coefficients based on wet-lab data; (e) revision of the acceptance criteria based on observed effect sizes. All such corrections will be documented in a Change Log addendum and version-controlled in the project repository.

The freeze statement does not preclude: (a) computational implementation of the framework (Phase 1 of the roadmap); (b) multi-expert AHP elicitation (Phase 2); (c) execution of the experimental programme (Phases 3-6); (d) manuscript drafting and submission (Phases 7-8). These activities are the execution of the frozen framework, not revisions to it.

Signed: Scientific Revision Committee. Date: as per document footer.